



**Note:**

- During the attendance check a sticker containing a unique code will be put on this exam.
- This code contains a unique number that associates this exam with your registration number.
- This number is printed both next to the code and to the signature field in the attendance check list.

## Mining Massive Datasets

**Exam:** IN2323 / Retake

**Date:** Monday 30<sup>th</sup> September, 2019

**Examiner:** Prof. Dr. Stephan Günnemann

**Time:** 08:00 – 09:30

	P 1	P 2	P 3	P 4	P 5	P 6	P 7	P 8
I								

### Working instructions

- This exam consists of **12 pages** with a total of **8 problems**.  
Please make sure that you received a complete copy of the exam.
- You can earn 43 points.
- **Detaching pages from the exam is prohibited!**
- Allowed resources:
  - A4 sheet of handwritten notes (two sides)
  - **no other materials (e.g. books, cell phones, calculators) are allowed!**
- Only write on the sheets given to you by supervisors. If you need more paper, ask the supervisors.
- Last two pages can be used as scratch paper.
- All sheets (including scratch paper) have to be returned at the end.
- **Only use a black or a blue pen (no pencils, red or green pens)!**
- Write your answers only in the provided solution boxes or the scratch paper.
- **For problems that say "Justify your answer" or "Show your work" you only get points if you provide a valid explanation.** Otherwise it's sufficient to only provide the correct answer.
- Exam duration - 90 minutes.

Left room from \_\_\_\_\_ to \_\_\_\_\_ / Early submission at \_\_\_\_\_

## Problem 1 AR models: stationarity (5 points)

Decide whether the following AR models are stationary or not. Everywhere  $\epsilon_t \sim \mathcal{N}(0, \sigma^2)$  with a positive  $\sigma$ .

Mark correct answers with a cross



To undo a cross, completely fill out the answer option



To re-mark an option, use a human-readable marking



a)  $\mathcal{X}_t = c + 0.1\mathcal{X}_{t-1} + \epsilon_t$  with some  $c \in \mathbb{R}$

☐ stationary for  $|c| > 0.1$ , otherwise non-stationary

☐ yes, always stationary

☐ no, always non-stationary

b)  $\mathcal{X}_t = -3 + 0.2\mathcal{X}_{t-1} - 0.01\mathcal{X}_{t-2} + c\epsilon_t$  with some  $c \in \mathbb{R} \setminus \{0\}$

☐ stationary for  $c > 0$ , otherwise non-stationary

☐ yes, always stationary

☐ no, always non-stationary

c)  $\mathcal{X}_t = 1 + 0.3\mathcal{X}_{t-1} - 0.03\mathcal{X}_{t-2} + 0.001\mathcal{X}_{t-3} + \epsilon_t$

☐ no, non-stationary

☐ yes, stationary

d)  $\mathcal{X}_t = -2 + 0.5\mathcal{X}_{t-n} + \epsilon_t$  with some  $n \in \mathbb{N}$

☐ no, always non-stationary

☐ stationary for  $n \leq 2$ , otherwise non-stationary

☐ yes, always stationary

e)  $\mathcal{X}_t = 2019 - \sum_{i=1}^n a^i \mathcal{X}_{t-i} + \epsilon_t$  with some  $n \in \mathbb{N}$ ,  $a \in \mathbb{R} \setminus \{0\}$

☐ stationary for  $|a| < 1$ , otherwise non-stationary

☐ stationary for  $|a|^n < 2019$ , otherwise non-stationary

☐ stationary for  $n = 1$ ,  $|a| < 1$ , otherwise non-stationary

## Problem 2 Hidden Markov Models (7 points)

Consider the following Hidden Markov Model where  $Z_t \in \{0, 1\}$  are latent variables and  $X_t \in \{a, b\}$  are discrete observed variables. We parametrize the prior and transition probabilities  $P(Z_1 = i) = \pi_i$ ,  $P(Z_{t+1} = j | Z_t = i) = A_{ij}$  and  $P(X_t = j | Z_t = i) = B_{ij}$  by:

$$\pi = \begin{bmatrix} 1/2 \\ 1/2 \end{bmatrix}, \quad A = \begin{bmatrix} \alpha & 1 - \alpha \\ 1/2 & 1/2 \end{bmatrix}, \quad B = \begin{bmatrix} 3/4 & 1/4 \\ 1/4 & 3/4 \end{bmatrix}$$

We assume we observed  $X = [a, b, a]$ .

a) Compute  $P(Z_2 | X_1, X_2)$  and  $P(Z_2 | X_1, X_2, X_3)$  as a function of  $\alpha$ .



0 ☐ b) What is the most probable state  $Z_2$  according to  $P(Z_2|X_1, X_2)$  and  $P(Z_2|X_1, X_2, X_3)$  for any value of  $\alpha \in [0, 1]$ .  
1 ☐ Hint:  $\sqrt{244} \approx 15.62$   
2 ☐

### Problem 3 RNNs & Word vectors (7 points)

You are solving a question-answering task. Given a context and a question, the goal is to find the answer inside the context. Bellow are two examples (1 and 2).

id	Context	Question	Answer
1	Mary was in the bathroom. Then she moved to the hallway.	Where is Mary?	hallway
2	John is in the hallway. Mary is there as well.	Where is Mary?	hallway

Assume that the question is represented with the vector  $\mathbf{q}$ . We want to know what is the probability that a word from the context is the answer. We decide to somehow represent every word  $w_i$  with an embedding  $\mathbf{h}_i$  and pass it together with  $\mathbf{q}$  through a neural network to get the probabilities. The only thing left to do is to decide how to get  $\mathbf{h}_i$ . We propose two approaches: sliding window and RNN.

a) **Sliding window** — Every word  $w_i$  is represented with a pretrained word vector  $\mathbf{v}_j$ . A sliding window of size 2 takes the neighbouring words and constructs the embedding for  $w_i$  as a sum of vectors:  $\mathbf{h}_i = \sum_{j=i-2}^{i+2} \mathbf{v}_j$ . Is it possible for this model to find the right answer in example 1? What about example 2? Justify.

☐ 0  
☐ 1  
☐ 2

b) **RNN** — As an alternative we use an RNN that takes pretrained word vectors  $\mathbf{v}_i$  from left to right and outputs  $\mathbf{h}_i$  as a word embedding. Why is this model able to output the right answer in example 1? Explain why it could fail answering example 2?

☐ 0  
☐ 1  
☐ 2

c) Propose another model that overcomes the shortcomings of the sliding window and the RNN. Describe what the input would be and how would you calculate the embedding  $\mathbf{h}_i$ . Explain why this model could work on both examples.

☐ 0  
☐ 1  
☐ 2  
☐ 3

## Problem 4 Deep Generative Model (5 points)

- 0 ☐ a) You are given a pseudo code implementation of 4 different variants of an AutoEncoder. Here,  $\mathbf{x}_i \in \mathbb{R}^d$   
 1 ☐ is the input data, and  $f_\theta : \mathbb{R}^d \mapsto \mathbb{R}^k$  and  $g_\phi : \mathbb{R}^k \mapsto \mathbb{R}^d$  are fully connected feed-forward neural networks  
 2 ☐ where  $\theta$  and  $\phi$  are learnable parameters. The final layers of  $f_\theta$  and  $g_\phi$  have no (i.e. have linear) activation  
 3 ☐ functions.  $\mathbf{I}_k$  is a  $k \times k$  identity matrix,  $\mathbf{0}_k$  is a  $k$ -dimensional vector of zeros,  $\mathcal{N}$  is the Normal distribution,  $\sigma$   
 is the softmax function, and  $\text{diag}(\mathbf{x})$  takes a vector  $\mathbf{x} \in \mathbb{R}^k$  and returns a  $k \times k$  diagonal matrix with the vector  
 values on the diagonal.

AutoEncoder 1

$$\begin{aligned}\epsilon_i &\sim \mathcal{N}(\mathbf{x}_i, \mathbf{I}_d) \\ \mathbf{h}_i &= f_\theta(\epsilon_i) \\ \tilde{\mathbf{x}}_i &= g_\phi(\mathbf{h}_i) \\ \mathcal{L} &= \sum_i \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|_2\end{aligned}$$

AutoEncoder 2

$$\begin{aligned}\mathbf{h}_i &= f_\theta(\mathbf{x}_i) \\ \epsilon_i &\sim \mathcal{N}(\mathbf{h}_i, \mathbf{I}_k) \\ \tilde{\mathbf{x}}_i &= g_\phi(\epsilon_i) \\ \mathcal{L} &= \sum_i \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|_2\end{aligned}$$

AutoEncoder 3

$$\begin{aligned}\mathbf{h}_i &= f_\theta(\mathbf{x}_i) \\ \epsilon_i &\sim \mathcal{N}(\mathbf{0}_k, \mathbf{I}_k) \\ \tilde{\mathbf{x}}_i &= g_\phi(\mathbf{h}_i + \epsilon_i) \\ \mathcal{L} &= \sum_i \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|_2\end{aligned}$$

Is it **necessary** to use the reparametrization trick in order to compute the gradients of  $\mathcal{L}$  w.r.t. to **both**  $\theta$  and  $\phi$  in the above implementations? Answer with Yes or No and provide a justification. If the answer is Yes, modify the pseudo code to implement the reparametrization trick.

AutoEncoder 1

AutoEncoder 2

AutoEncoder 3

b) Assume the same setup as in a). The model specified by the following pseudo code is **not well defined**. Specify the reason why, and modify the pseudo code such that the model becomes well-defined. In addition, if you think it is **necessary** to use the reparametrization trick, please include it in your implementation.

☐ 0  
☐ 1  
☐ 2

$$\begin{aligned}\mathbf{h}_i &= f_{\theta}(\mathbf{x}_i) \\ \epsilon_i &\sim \mathcal{N}(\mathbf{0}_k, \text{diag}(\mathbf{h}_i)) \\ \tilde{\mathbf{x}}_i &= g_{\phi}(\epsilon_i) \\ \mathcal{L} &= \sum_i \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|_2\end{aligned}$$

## Problem 5 Spectral clustering (3 points)

0 ☐  
1 ☐  
2 ☐  
3 ☐

Given is the following matrix  $M \in \mathbb{R}^{9 \times 9}$ .

$$M = \begin{bmatrix} 1 & 0 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & -1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 2 & -1 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 2 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & -1 & 3 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & -1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & -1 & 1 \end{bmatrix}$$

Write down the **exact** value of the three smallest eigenvalues  $\lambda_1, \lambda_2, \lambda_3$  of  $M$  and the respective eigenvectors  $x_1, x_2, x_3$ .

## Problem 6 Spectral clustering (3 points)

0 ☐  
1 ☐  
2 ☐  
3 ☐

You are given an undirected graph  $G = (V, E)$ . It is known that the second smallest eigenvalue of the unnormalized Laplacian  $L = D - W$  is equal to 10.

Let  $\phi(G)$  denote the best possible ratio cut achievable on the graph  $G$

$$\phi(G) = \min_{S \subset V} \text{ratio-cut}(S, \bar{S})$$

What are possible values of  $\phi(G)$  for the given graph? Select all that apply. Justify your answer.

- a) 1
- b) 2
- c) 4
- d) 8
- e) 16



## Problem 7 Ranking (5 points)

Given a graph  $G$  with 5 nodes, assume that you have access to several topic-sensitive PageRank vectors, each pre-computed using a **different** teleport set  $S$ , and the **same (fixed)** teleport parameter  $\beta$ ,  $0 < \beta < 1$ .

- $\pi_{235} \in \mathbb{R}^5$ , with teleport set  $S = \{2, 3, 5\}$
- $\pi_{124} \in \mathbb{R}^5$ , with teleport set  $S = \{1, 2, 4\}$
- $\pi_{134} \in \mathbb{R}^5$ , with teleport set  $S = \{1, 3, 4\}$
- $\pi_3 \in \mathbb{R}^5$ , with teleport set  $S = \{3\}$

Assume that the random walker always teleports **uniformly at random** to each node in the teleport set.

Is it possible to compute each of the following PageRank vectors without access to the graph  $G$ , i.e. using only the above pre-computed vectors? If so, specify the exact equation as a function of  $\pi_{235}$ ,  $\pi_{124}$ ,  $\pi_{134}$  and  $\pi_3$ . If not, justify why not.

a) Is it possible to compute  $\pi_{14} \in \mathbb{R}^5$  with teleport set  $S = \{1, 4\}$ ?

☐ 0  
☐ 1

b) Is it possible to compute  $\pi_5 \in \mathbb{R}^5$  with teleport set  $S = \{5\}$ ?

☐ 0  
☐ 1

c) Is it possible to compute  $\pi_1 \in \mathbb{R}^5$  with teleport set  $S = \{1\}$ ?

☐ 0  
☐ 1

d) Is it possible to compute  $\pi_w \in \mathbb{R}^5$  with teleport set  $S = \{1, 2, 3, 4, 5\}$ , where we do not teleport to each node uniformly at random but rather with weights 0.2, 0.3, 0.1, 0.2, 0.2, respectively?

☐ 0  
☐ 1  
☐ 2

## Problem 8 Graph Neural Networks (8 points)

Given an unweighted, undirected graph  $G$  with adjacency matrix  $\mathbf{A} \in \{0, 1\}^{N \times N}$  and node attribute matrix  $\mathbf{X} \in \mathbb{R}^{N \times D}$ , your task is to perform semi-supervised node classification with  $C$  classes using a graph convolutional network (GCN).

In matrix notation, a GCN with  $K$  layers is recursively defined as follows.

$$\begin{aligned} \mathbf{H}^{(0)} &= \mathbf{X} \\ \mathbf{H}^{(k)} &= \sigma^{(k)} \left( \tilde{\mathbf{A}} \mathbf{H}^{(k-1)} \mathbf{W}^{(k)} \right) \quad \text{for } k \in \{1, \dots, K\} \end{aligned}$$

That is,  $\mathbf{H}^{(K)} \in \mathbb{R}^{N \times C}$  contains the class predictions for **all nodes** stacked in a matrix. Here,  $\sigma^{(k)}$  is the  $\text{ReLU}(\cdot)$  activation function for  $k \in \{1, \dots, K-1\}$  and the  $\text{softmax}(\cdot)$  function for the final (output) layer  $k = K$ .  $\mathbf{W}^{(k)}$  is the weight matrix of layer  $k$ .

$\tilde{\mathbf{A}} \in \mathbb{R}^{N \times N}$  is a degree-normalized version of the adjacency matrix whose entries are

$$\tilde{\mathbf{A}}_{uv} = \begin{cases} \frac{1}{\sqrt{d_u d_v}} & \text{if } \mathbf{A}_{uv} = 1 \\ \frac{1}{d_u} & \text{if } u = v \\ 0 & \text{else,} \end{cases}$$

where  $d_u$  is the degree of node  $u$ .

- 0 ☐ a) GCN is an instance of the differentiable message passing framework. Provide the corresponding message function  $M$  and update function  $U$  of GCN. For simplicity, you may write  $\sigma^{(k)}$  for all layers.

1 ☐

2 ☐

3 ☐

- 0 ☐ b) Assume you have a GCN with 3 layers, i.e.  $K = 3$ . Provide the non-recursive (unrolled) expression for  $\mathbf{H}^{(3)}$ . That is, write down the single-line equation of  $\mathbf{H}^{(3)}$  in matrix form.

1 ☐

☐ 0
 ☐ 1
 ☐ 2

c) We consider now the two following situations:

- (1)  $\mathbf{A}_{uv} = 1$  for  $u = v$  and 0 else.
- (2)  $\sigma^{(k)}(x) = x$ , i.e. identity activation function, for  $k \in \{1, 2\}$

Apart from the additional information in each situation, the models are identical to the GCN defined above. Simplify the single-line, matrix-form expression of  $\mathbf{H}^{(3)}$  given the additional information provided in each situation.

d) For each situation, find one equivalent model in the table below. **You may select each option only once.** Briefly justify your answer for each situation.

☐ 0
 ☐ 1
 ☐ 2

Recurrent neural network (RNN)	Linear regression
Feed-forward neural network (FFNN)	Label Propagation (LP)
Linear function	Deep Generative Model
Logistic regression on $\mathbf{X}$	Logistic regression on pre-processed features $\tilde{\mathbf{X}}$

**Additional space for solutions—clearly mark the (sub)problem your answers are related to and strike out invalid solutions.**

