

Machine Learning for Graphs and Sequential Data

Robustness of Machine Learning Models

Lecturer: Prof. Dr. Stephan Günnemann

www.daml.in.tum.de

Summer Term 2020

Data Analytics and
Machine Learning 

Roadmap

1. **Introduction**
2. Construction of adversarial examples
3. Improving robustness
4. Certifiable robustness
 - Exact certification
 - Convex relaxations
 - Lipschitz-continuity
 - Randomized smoothing

Introduction

- Often ML models and algorithms are optimized w.r.t. simple metrics
 - e.g. misclassification rate, reconstruction error, etc.
- As ML/AI is becoming more widespread and is used in critical applications (e.g. autonomous driving, algorithmic decision-making involving humans) we must consider further aspects
- As ML models get deployed in the real-world they create feedback loops which can have potentially unintended consequences
- One important aspect: Are the ML models reliable?
 - How do they behave in the wild? When your data might, e.g., be corrupted?

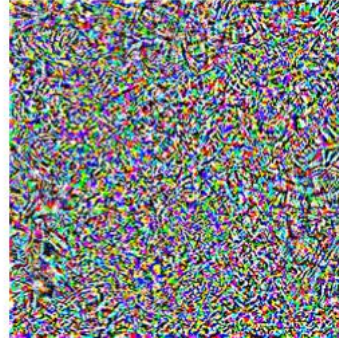
What Are Adversarial Examples?

Predicted
class:

“pig”



+ 0.005 x



=

?



Image from http://gradientscience.org/intro_adversarial/

What Are Adversarial Examples?

Predicted
class:


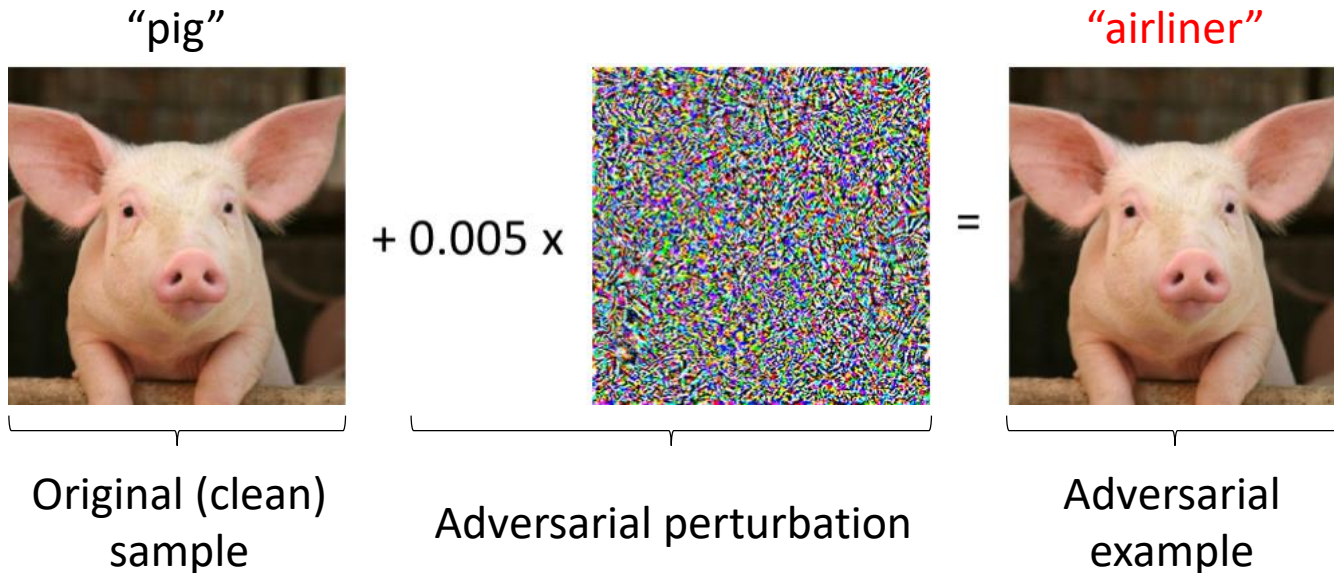


Image from http://gradientscience.org/intro_adversarial/

What Are Adversarial Examples?

Predicted
class:



Small (imperceptible) but **specifically crafted perturbations** lead to **false predictions** in machine learning models.

- Why should we care about adversarial examples?
- What does “small” mean?
- How are adversarial examples created?
- How to protect against adversarial examples?

Image from http://gradientscience.org/intro_adversarial/

Why We (Should) Care About Adversarial Examples

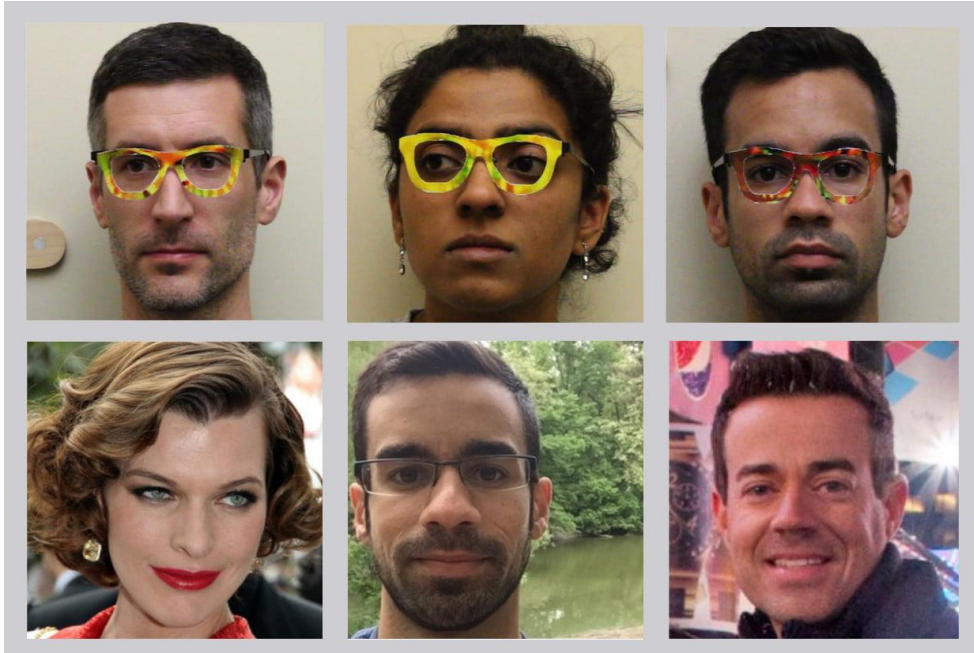
Real-world risks

- Adversarial examples are an obvious security threat for many real-world applications, e.g. self-driving cars.
- Adversarial examples also exist in the real world, e.g. 2D / 3D prints, special glasses to disturb face recognition, etc.

Conceptual gaps

- Neural networks are hypothesized to learn meaningful representations that capture semantic understanding of the domain and task.
 - Adversarial examples are counterexamples to this hypothesis: the semantic content of the samples is unchanged but the network is fooled.
- Nature as an adversary: Even if there is no adversary in our use-case, we should quantify robustness to **worst-case noise**

Adversarial Examples



Adversarial glasses fool facial recognition systems into classifying the wearer as someone else, [Sharif et al., 2016]



ML systems classify the adversarially modified STOP sign as a speed limit sign, [Eykholt et al., 2018]

Adversarial Examples – Definition

Classification task:

- Dataset: $(\mathbf{x}_i, y_i) \sim \mathbb{P}_{\text{data}}, \quad (\mathbf{x}_i, y_i) \in \mathbb{R}^d \times \mathcal{Y}$
- Classifier: $f: \mathbb{R}^d \rightarrow \mathcal{Y}$

- Specify a perturbation set $\mathcal{P}(\mathbf{x})$, i.e. a set of perturbations which when applied to \mathbf{x} do not change it's semantic
 - and, thus, should also not change it's classification
- We say that a point $\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})$ is an adversarial example for f at (\mathbf{x}, y)
 - if $f(\mathbf{x}) = y$, i.e., f correctly classifies \mathbf{x}
 - but $f(\tilde{\mathbf{x}}) \neq y$, i.e. fails to correctly classify $\tilde{\mathbf{x}}$

On “small” Perturbations

- Perturbations should not change the **semantic content** of a sample.
- This is often translated into L_p constraints with some small ϵ
- L_p norm: $\mathcal{P}_{\epsilon,p}(\mathbf{x}) = \{\tilde{\mathbf{x}}: \|\tilde{\mathbf{x}} - \mathbf{x}\|_p < \epsilon\}$, typically $p \in \{1, 2, \infty\}$
- While mathematically convenient, L_p norms with small ϵ do not contain **all** semantically meaningless perturbations.
- For example, a **small rotation** does typically not change the meaning of a picture but often corresponds to **large changes** in L_p norm.

Attack Variants

- **Evasion attacks:** given a **fixed**, trained classifier f , the attacker aims to find an adversarial perturbation (at test time)
- **Poisoning attacks:** the adversary aims to modify the **training dataset** such that a classifier trained on the dataset has properties desired by the attacker.
 - i.e. the manipulation/corruption is done **before** training
 - ➔ **Not covered in this course.**
- **Targeted attacks:** the attacker aims to have a certain sample classified as a specific class (e.g. speed limit 100 km/h sign).
- **Untargeted attacks:** the attacker aims to have a sample misclassified as **any** class different than the correct one.

Roadmap

1. Introduction
- 2. Construction of adversarial examples**
3. Improving robustness
4. Certifiable robustness
 - Exact certification
 - Convex relaxations
 - Lipschitz-continuity
 - Randomized smoothing

Adversarial Attacks: Objective Function

Construction of adversarial examples can be phrased as an optimization problem. For example:

$$\tilde{\mathbf{x}}_{\mathbf{x}}^* = \arg \max_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell_{0/1}(f(\tilde{\mathbf{x}}), y)$$

- **Recall:** $\ell_{0/1}$ is the zero/one loss (0 if correct, 1 if incorrect).
- **However:** $\ell_{0/1}$ has either zero or undefined gradient.
- Therefore, the **cross-entropy loss** \mathcal{L} is often used as a surrogate:

$$\tilde{\mathbf{x}}_{\mathbf{x}}^* = \arg \max_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \mathcal{L}(f(\tilde{\mathbf{x}}), y)$$

Projected Gradient Descent

$$\tilde{\mathbf{x}}_{\mathbf{x}}^* = \arg \max_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \mathcal{L}(f(\tilde{\mathbf{x}}), y)$$

- One common method is **Projected gradient descent** (PGD): after each gradient step on the objective, project onto the valid domain.
- $\mathbf{x}_{t+1} = \Pi(\mathbf{x}_t + \eta_t \nabla_{\mathbf{x}} \mathcal{L}(f(\mathbf{x}_t), y))$
- Like training the model but **updating the data instead of the weights**.
- Note: since $f(\tilde{\mathbf{x}})$ is **not convex**, in general we cannot find the global optimum.

Fast Gradient-Sign Method (FGSM):

- $\tilde{\mathbf{x}} = \Pi(\mathbf{x} + \eta \cdot \text{sign}(\nabla_{\mathbf{x}} \mathcal{L}(f(\mathbf{x}), y)))$
- When $\mathcal{P}(\mathbf{x})$ is a ball with radius ϵ measured by the L_{∞} norm, setting $\eta = \epsilon$ yields valid perturbations with only a single step and without projection.

Alternative Optimization Problem

- An alternative formulation is to optimize:

$$\min_{\tilde{\mathbf{x}}} \mathcal{D}(\mathbf{x}, \tilde{\mathbf{x}}) \quad \text{subject to } \ell_{0/1}(f(\tilde{\mathbf{x}}), y) > 0$$

- Here, \mathcal{D} is a term that is large when $\tilde{\mathbf{x}}$ is far from \mathbf{x} (e.g., an L_p distance)
- [Carlini and Wagner, 2017] convert this constrained into an **unconstrained** optimization problem:

$$\min_{\tilde{\mathbf{x}}} \mathcal{D}(\mathbf{x}, \tilde{\mathbf{x}}) + \lambda \cdot L(\tilde{\mathbf{x}}, y),$$

- A very effective **loss function** is

$$L(\tilde{\mathbf{x}}, y) = \left[Z(\tilde{\mathbf{x}})_y - \max_{i \neq y} (Z(\tilde{\mathbf{x}})_i) \right]_+$$

- y is the original class we want $\tilde{\mathbf{x}}$ to deviate from
- $[\mathbf{x}]_+ = \max(\mathbf{x}, 0)$
- $Z(\tilde{\mathbf{x}})_i = \log f(\tilde{\mathbf{x}})_i$ (log probability of class i)
- The loss L is positive if $\tilde{\mathbf{x}}$ is classified as y and 0 otherwise.

Roadmap

1. Introduction
2. Construction of adversarial examples
- 3. Improving robustness**
4. Certifiable robustness
 - Exact certification
 - Convex relaxations
 - Lipschitz-continuity
 - Randomized smoothing

Introduction

- Most ML models (trained and/or applied in the traditional way) are vulnerable to adversarial examples
- How to defend against adversarial examples?
- Can we prevent them?
- Can we improve the robustness of our models?

What does not seem to work

- **Post-hoc prevention of attacks**

E.g. gradient obfuscation, i.e. randomizing or shattering gradients of the model in order to prevent gradient-based attacks. So far, these defenses have all been broken by stronger attacks.

- **Detection of adversarial examples**

E.g. out-of-distribution shift: since the adversarial examples come from a different distribution as the natural images, we could try to distinguish the two distributions or perform outlier / anomaly detection. While some of these methods work against “vanilla” PGD attacks, targeted attacks are very successful against these defenses.

- Fixing a “bad” model seems not to be the solution

Robust Training

Robust training refers to training procedures aimed at producing models that are robust to adversarial (and/or other) perturbations.

A **common theme** is to optimize a ‘worst-case’ loss (also called robust loss), i.e. the loss achieved under the worst-case perturbation.

- Let $\ell(\hat{y}, y)$ be some loss, e.g. cross-entropy loss
- The (non-robust) training tries to find an f that minimizes the expected loss

$$R = \mathbb{E}_{(\mathbf{x}, y) \in \mathbb{P}_{\text{data}}} [\ell(f(\mathbf{x}), y)]$$

- The robust version of this problem is

$$R_{\text{rob}} = \mathbb{E}_{(\mathbf{x}, y) \in \mathbb{P}_{\text{data}}} \left[\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f(\tilde{\mathbf{x}}), y) \right]$$

Loss achieved by the worst-case perturbation in $\mathcal{P}(\mathbf{x})$

Robust Training

Robust training refers to training procedures aimed at producing models that are robust to adversarial (and/or other) perturbations.

A **common theme** is to optimize a ‘worst-case’ loss (also called robust loss), i.e. the loss achieved under the worst-case perturbation.

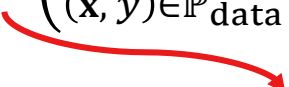
- **Adversarial training** is an easy-to-implement robust training procedure that uses adversarial examples as a proxy for the ‘worst-case’ perturbation.
- In the next chapter, we will cover **robustness certification** techniques. Some of these can also be used for robust training.

Adversarial Training

Idea: perform stochastic gradient descent (SGD) on the **robust loss** R_{rob} :

$$R_{\text{rob}} = \mathbb{E}_{(\mathbf{x}, y) \in \mathbb{P}_{\text{data}}} \left[\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f(\tilde{\mathbf{x}}), y) \right]$$

For $f = f_{\theta}$ being a **neural network** parameterized by weights θ , we can write

$$\begin{aligned} \nabla_{\theta} R_{\text{rob}} &= \nabla_{\theta} \left(\mathbb{E}_{(\mathbf{x}, y) \in \mathbb{P}_{\text{data}}} \left[\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f_{\theta}(\tilde{\mathbf{x}}), y) \right] \right) \\ &= \mathbb{E}_{(\mathbf{x}, y) \in \mathbb{P}_{\text{data}}} \left[\nabla_{\theta} \left(\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f_{\theta}(\tilde{\mathbf{x}}), y) \right) \right] \end{aligned}$$


How to take the gradient of the **worst-case loss** w.r.t. the weights θ ?

Adversarial Training: Danskin's Theorem

- How to obtain $\nabla_{\theta} L = \nabla_{\theta} \left(\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f_{\theta}(\tilde{\mathbf{x}}), y) \right)$?
- That is, the **gradient** of the **worst-case loss** w.r.t. the model parameters.
- **Danskin's Theorem***: Let $\Delta(\theta)$ be the set of $\tilde{\mathbf{x}}$ for which the supremum is obtained. If $\Delta(\theta)$ contains only a single element, i.e. $\Delta(\theta) = \{\tilde{\mathbf{x}}_{\theta}^*\}$, then the sup is differentiable at θ and

$$\nabla_{\theta} \left(\sup_{\tilde{\mathbf{x}} \in \mathcal{P}(\mathbf{x})} \ell(f_{\theta}(\tilde{\mathbf{x}}), y) \right) = \nabla_{\theta} \ell(f_{\theta}(\tilde{\mathbf{x}}_{\theta}^*), y)$$

* Technically, the theorem requires some conditions which might not hold in our case; e.g., that $\ell(f_{\theta}(\tilde{\mathbf{x}}), y)$ is convex in θ .

Adversarial Training: Algorithm

- Using **Danskin's theorem** we can compute the gradient of the worst-case loss given the corresponding perturbation.
- **Problem:** finding the **worst-case** perturbed example $\tilde{\mathbf{x}}$ is **intractable**; if we could find it efficiently, we would have solved the **exact verification** problem.
- **Idea:** Create any adversarial example as a proxy of the worst-case perturbation, e.g. via the fast gradient-sign method (FGSM).
- Adversarial training algorithm outline:
 1. Sample $(\mathbf{x}_i, y_i) \sim \mathbb{P}_{\text{data}}$
 2. Using an adversarial attack procedure, find an $\tilde{\mathbf{x}}_i$ with high loss $\ell(f_{\theta}(\tilde{\mathbf{x}}_i), y_i)$
 3. Update weights via gradient descent: $\theta \leftarrow \theta - \eta \nabla_{\theta} \ell(f_{\theta}(\tilde{\mathbf{x}}_i), y_i)$
- In step 2, we must trade off the strength of the attack with its computational cost.

Adversarial Training: Summary

Pro:

- It empirically increases robustness of the resulting models.
- It is easy to implement.

Con:

- If we want to use a powerful attack on the inner optimization, the slowdown is about 10x compared to standard training.
- The resulting models typically have lower accuracy on clean data.
- We don't get any theoretical guarantees of the model's robustness

Questions – Rob1

1. Given an arbitrary binary classifier f for an input domain \mathbb{R}^d and the perturbation set $\mathcal{P}_{\epsilon,p}(\mathbf{x})$ as defined before. Is it possible that every $\mathbf{x} \in \mathbb{R}^d$ is “robust”, i.e. no adversarial example exists?
2. Will the fast gradient-sign method (FGSM) always find an adversarial example (assuming there exist some in the set of perturbations $\mathcal{P}_{\epsilon,\infty}(\mathbf{x})$)?

Recommended Reading

- Lecture 09: Introduction to adversarial examples and
Lecture 10: Empirical defenses for adversarial examples
of Jerry Li's course on Robustness in Machine Learning (CSE 599-M),
<https://jerryzli.github.io/robust-ml-fall19.html>

References

- Carlini, Nicholas, and David Wagner. "Towards evaluating the robustness of neural networks." *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2017.
- Eykholt, Kevin, et al. "Robust physical-world attacks on deep learning visual classification." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2018.
- Sharif, Mahmood, et al. "Accessorize to a crime: Real and stealthy attacks on state-of-the-art face recognition." *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*. 2016.