

Machine Learning for Graphs and Sequential Data

Graphs – Graphs & Networks

Lecturer: Prof. Dr. Stephan Günnemann

www.daml.in.tum.de

Summer Term 2020

Data Analytics and
Machine Learning 

Roadmap

- **Chapter: Graphs**

- 1. Graphs & Networks**

- **Motivation & Definitions**
 - Properties of Real Networks

- 2. Generative Models

- 3. Clustering



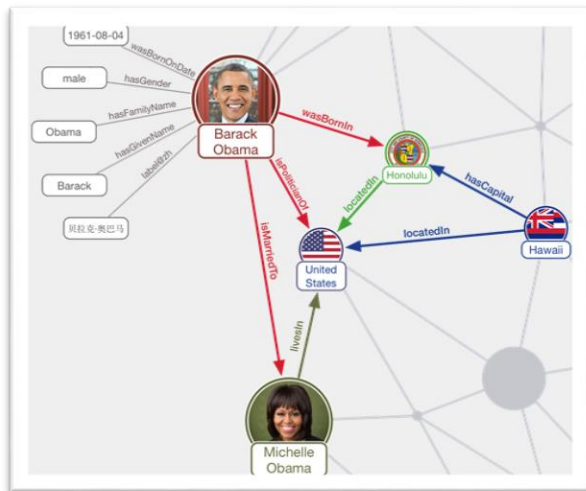
- 4. Node Embeddings

- 5. Ranking

- 6. Semi-Supervised Learning

- 7. Limitations of GNNs

Why Should We Care?



- Web: hyper-text graph
- Information retrieval: bi-partite graphs (documents-terms)
- E-commerce: User-Product graphs

Why Should We Care?

- ‘Viral’ marketing, News propagation
- Computer network security: email/IP traffic and anomaly detection
- Ranking of search results
- Fraud detection in e-commerce systems
- Drug discovery, molecule property prediction
- Scene graph analysis
- ...

Basic Definition

- Plain/simple graph $G = (V, E)$
 - Set of nodes V
 - Set of edges $E \subseteq V \times V$ // for undirected graphs: $(i, j) \in E \Leftrightarrow (j, i) \in E$
- Equivalent representation via (binary) adjacency matrix $A \in \{0,1\}^{|V| \times |V|}$
- Multiple extensions possible
 - weighted graphs (node weights, edge weights)
 - attributed graphs (multi-dimensional vectors assigned to nodes/edges)
 - temporal graphs (timestamp associated with node/edge)

Roadmap

- **Chapter: Graphs**

- 1. Graphs & Networks**

- Motivation & Definitions
 - **Properties of Real Networks**

- 2. Generative Models

- 3. Clustering

- 4. Node Embeddings



- 5. Ranking

- 6. Semi-Supervised Learning

- 7. Limitations of GNNs

How are real networks structured?


- What does the Internet look like?
- What does Facebook look like?
- What is 'normal'/'abnormal'?
- Which patterns/laws hold?
 - To spot anomalies (rarities), we have to discover patterns
 - **Large datasets** reveal patterns/anomalies that may be invisible otherwise...



Are real graphs random?


- Erdős-Renyi Random Graph Model

- Start with N (isolated) nodes
- For every pair $v_1, v_2 \in V$ add an edge with probability p
- Every edge occurs with equal probability




- Example: 100 nodes, avg degree = 2

before layout



after layout




- No obvious patterns

Laws and Patterns

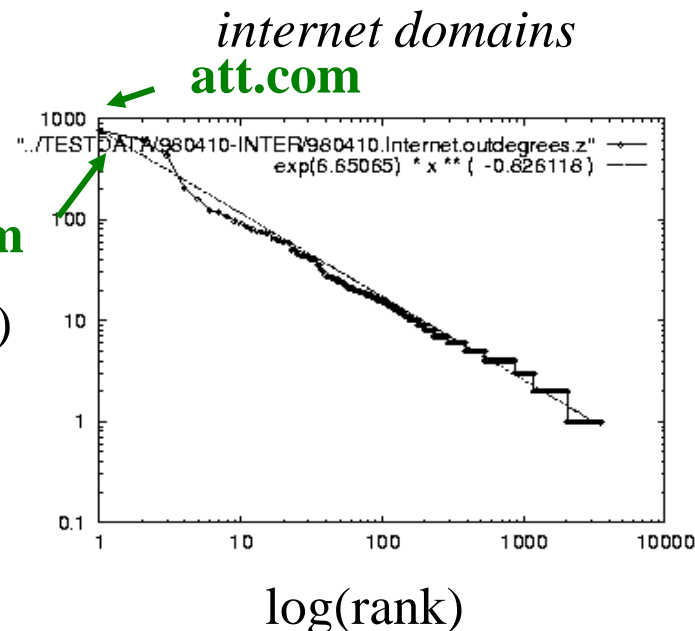
- Q: Are real graphs random?
- A: NO!
 - Diameter
 - in- and out- degree distributions
 - other (surprising) patterns
- So, let's look at the data

Power Law Distributions

- Gaussian distributions are common in nature
- In networks, however, a **power law distribution** often explains the data better
- Example: Power law in the degree distribution




ibm.com
log(degree)



Power Law Distributions

- Definition Power Law:
 - Two variables x and y are related by a power law when $y(x) = Ax^{-\gamma}$ where A and γ (power law exponent) are positive constants
 - A random variable is distributed according to a power law when the probability density function (pdf) is given by $p(x) = Ax^{-\gamma}$ with $\gamma > 1$

- Note: Power law distribution looks like a line on a log-log scale
- **Characteristic: Decay of pdf is only polynomial (Gaussian: exponential)**
- → More likely to observe values far to the right of the mean

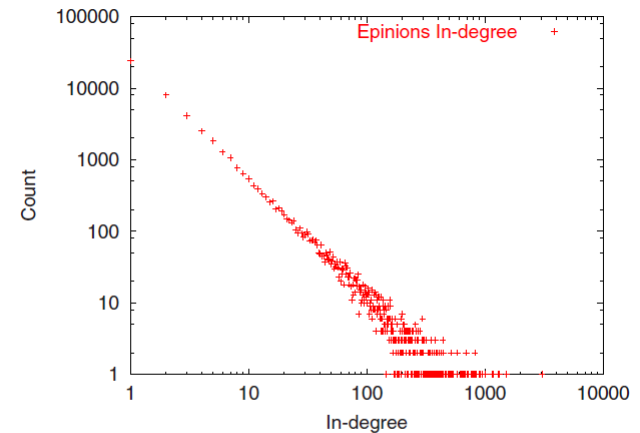


a.k.a. heavy-tailed distributions

- E.g. more likely to have nodes with a very high degree

Examples: Power Law Distributions

- "Internet AS" graph with exponent 2.1 – 2.2
- "Internet router" graph with exponent 2.48
- Citation graphs with exponent 3
- Epinions (who-trust-whom)
- ...

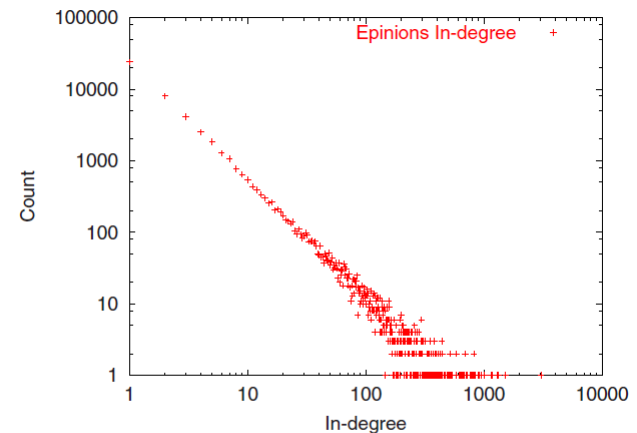


(a) Epinions In-degree

- Note: Graphs with degree distributions following a power law are called **scale-free**
 - $y(a \cdot x) = b \cdot y(x)$ // $y(x)$ =number of nodes with degree x

Important Remark


- Important: We do not claim that the data's underlying distribution is a power law distribution
- Instead: In many cases a power law distribution is a **good description/fit** (or approximation) of what we observe
 - usually deviations to power law distribution are observed (to different degrees)
 - other models: exponential cutoff, lognormal distributions



(a) Epinions In-degree

"Gaussian Trap"

- Q: So what?
- A1: Be careful when writing algorithms!
 - Example: # of two-step-away pairs (= friends of friends)
 - $O(d_{\max}^2) \sim 10M^2$
for storage: $\sim 0.8\text{PB} \rightarrow$ a data center(!)




Patterns and Algorithms

- Q: So what?
- A2: Patterns allow to design new algorithms!

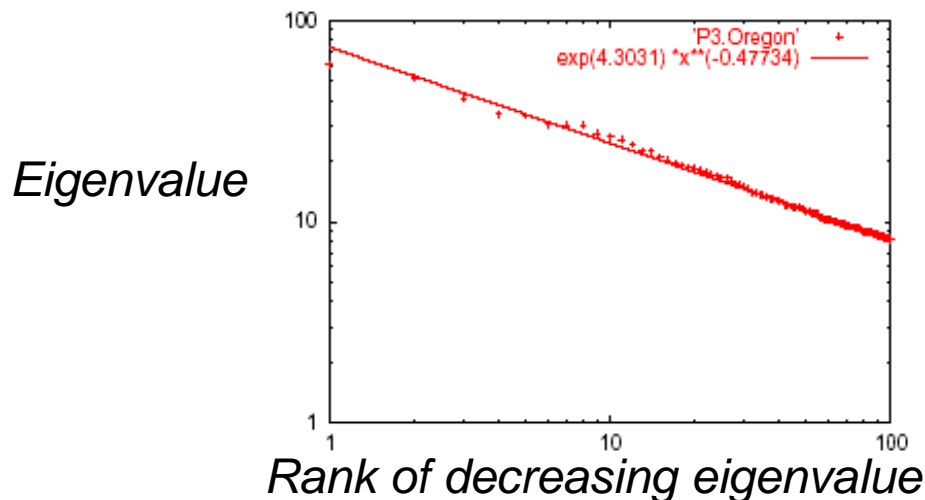
Example: Clustering Coefficient

- Real social networks have a lot of triangles
 - Friends of friends are friends
- Clustering Coefficient
 - $C = \frac{3 \cdot \text{number of triangles}}{\text{number of connected triplets}}$
 - Real-world networks often show large clustering coefficients
→ Strong community structure
- Computation of clustering coefficient requires computation of triangles
 - But: **triangles are expensive to compute** (3-way join)



Efficient Triangle Counting

- How to efficiently estimate the number of triangles?
 - Let's look at some patterns...
- Recap: Adjacency matrix $A \in \{0, 1\}^{N \times N}$
 - $a_{ij} = 1$ if there is an edge between i and j , 0 otherwise
- Recap: Eigenvalue decomposition $A x = \lambda x$
- Observer: Power law in the eigenvalues of the adjacency matrix



Exponent = slope

$$E = -0.48$$

Efficient Triangle Counting

- How does this help?
- 1. Some nice fact (easy to show):
 - number of triangles $= \frac{1}{6} \text{trace}(A^3) = \frac{1}{6} \sum_i \lambda_i^3$
 - λ_i = eigenvalues of adjacency matrix A
- 2. Eigenvalues follow power law (highly skewed)
 - we only need the top few (largest) eigenvalues!
 - how can we compute them efficiently?

Recap: Power Iteration

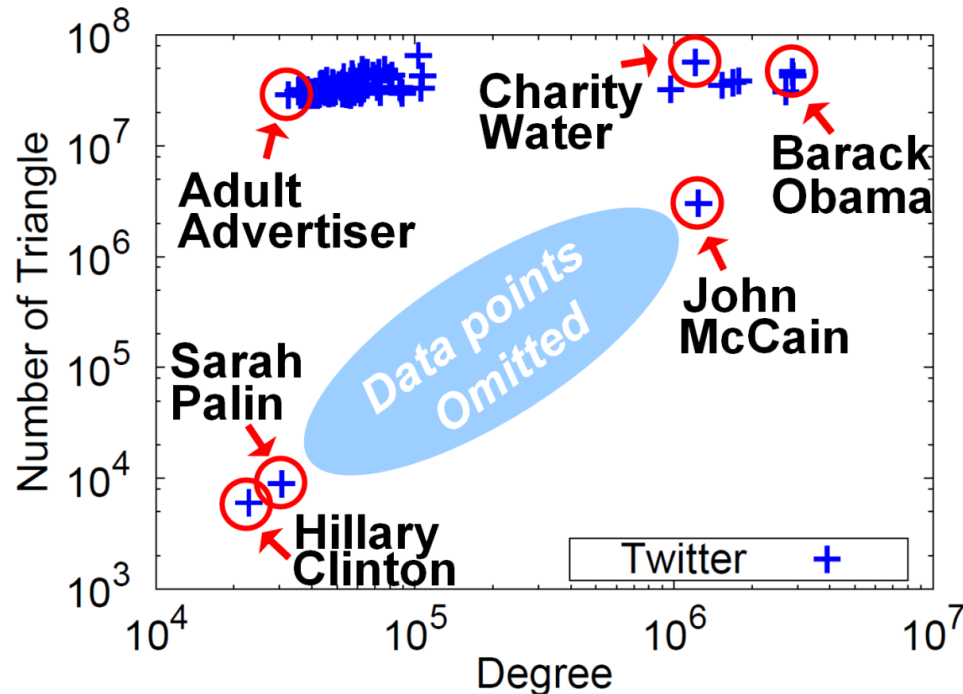
- Eigenvalues are important for many ML/data mining tasks
 - PCA, Ranking of Websites, Community Detection, ...
 - How to compute them efficiently?
- Power iteration (a.k.a. Von Mises iteration)
 - Iterative approach to compute **a single** eigenvector
- Let X be a matrix and v be an arbitrary (normalized) vector
 - Iteratively compute $v \leftarrow \frac{X \cdot v}{\|X \cdot v\|}$ until convergence
 - in each step, v is simply multiplied with X and normalized
 - **v converges to the eigenvector of X with greatest absolute value**
 - Highly efficient for sparse data

Recap: Power Iteration

- Convergence:
 - Linear convergence with rate $|\lambda_2/\lambda_1|$
 - Fast convergence if first and second eigenvalue are dissimilar

- How to find **multiple (the k largest) eigenvectors**?
 - Let us focus on symmetric matrices X
 - Eigenvalue decomposition leads to: $X = \Gamma \cdot \Lambda \cdot \Gamma^T = \sum_{i=1}^d \lambda_i \cdot \gamma_i \cdot \gamma_i^T$
 - Define deflated matrix: $\hat{X} = X - \lambda_1 \cdot \gamma_1 \cdot \gamma_1^T$
 - \hat{X} has the same eigenvectors as X except the first one has become zero
 - Apply power iteration on \hat{X} to find the second largest eigenvector of X


Example: Analysis of Twitter



- Anomalous nodes in Twitter (~ 3 billion edges)
- [U Kang et. al., PAKDD'11]

Real Graphs are Sparse

- N^2 possible edges for a graph with N nodes
- However, real-world graphs are very sparse $E \ll N^2$
- Instead of $E = O(N^2)$, we see $E = O(N^\alpha)$ with α significantly less than 2



Every “XX” is a real world network

Note the log-log scale

$$\alpha \approx 1.4$$

Small World Phenomenon


- Famous experiment by Travers and Milgram [TM1969]
 - Setup: Try to reach a random person by sending a chain letter
 - Result: The average length of chains that reach the person was six
 - → Length very small compared to number of participants
 - → „Small world phenomenon“; „six degrees of separation“
- Ways to measure this phenomenon
 - Characteristic path length
 - Average diameter
 - Effective diameter/Eccentricity



Characteristic Path Length & Average Diameter

■ Characteristic path length

- For each starting node $v \in V$ consider the shortest path to every other node
- Take the average length of all these paths
- Consider average path length for all starting nodes and take the median
- $\text{median}_{v \in V} \left\{ \frac{1}{|V|} \sum_{v_j \in V} \text{len}(p_{\min}(v, v_j)) \right\}$




■ Average diameter

- Similar to above but use (in last step) the mean instead of the median
- $\frac{1}{|V|} \sum_{v \in V} \frac{1}{|V|} \sum_{v_j \in V} \text{len}(p_{\min}(v, v_j))$

Effective Diameter / Eccentricity & Hop-plot

- Let $N_h(u)$ be the number of nodes reachable from u via h hops
 - $N_h(u) = \{v \in V \mid \text{len}(p_{\min}(u, v)) \leq h\}$
- The total neighborhood size N_h is the sum over all starting nodes
 - $N_h = \sum_{u \in V} |N_h(u)|$
- Hop-plot: Plot of N_h versus h
- Effective diameter (or Eccentricity)
 - Minimum number of hops in which some fraction (e.g. 90%) of all connected pairs of nodes can reach each other
 - $\min\{k \in \mathbb{N} \mid N_k \geq 0.9 \cdot |V|^2\}$
 - **Advantage: Also works for disconnected graphs**



Importance of „Network Laws“

- Laws describing „normal“ networks are important for:
- Design of algorithms
- Detection of abnormal/interesting patterns
 - Abnormalities deviate from the „normal“ patterns
 - Prerequisite: specify what is normal
- Development of graph generators
 - Often: real world data not public available; or just small excerpts
 - Use synthetic data to test algorithms
 - Requirement: generate synthetic but realistic graphs
- Simulation studies
 - E.g. test next-generation internet protocol on graph „similar“ to what Internet will look like a few years into the future

Questions

- How much memory do you need to store the edges of a graph with 1000 nodes and 10,000 edges in a dense adjacency matrix? How much for a sparse matrix?
- What is the average degree in an Erdős-Renyi graph with edge probability p ? And in a real world sparse graph with $O(E) = O(N^\alpha)$?

Reading Material

- "Graph Mining: Laws, Tools, and Case Studies" by Deepayan Chakrabarti, Christos Faloutsos