

Machine Learning for Graphs and Sequential Data

Graphs – Ranking

Lecturer: Prof. Dr. Stephan Günnemann

www.daml.in.tum.de

Summer Term 2020

Data Analytics and
Machine Learning 


Roadmap

- **Chapter: Graphs**

1. Graphs & Networks
2. Generative Models
3. Clustering
4. Node Embeddings
- 5. Ranking**
6. Semi-Supervised Learning
7. Limitations of GNNs

Motivation: Ranking of Nodes

- How to organize the Web?
- First try: Human curated Web directories
 - Yahoo, DMOZ, LookSmart
- Second try: Web Search
 - Information Retrieval investigates: Find relevant docs in a small and trusted set
 - Newspaper articles, Patents, etc.
 - But: **Web is huge**, full of untrusted documents, randomness, web spam, etc.
- Web pages are not equally “important”
 - `www.some-personal-website.com` vs. `www.tum.de`
- There is large diversity in the web-graph node connectivity.
Let's rank the pages by the link structure!




PageRank

- Core idea: **A page is important if many important pages point to it**
 - recursive formulation

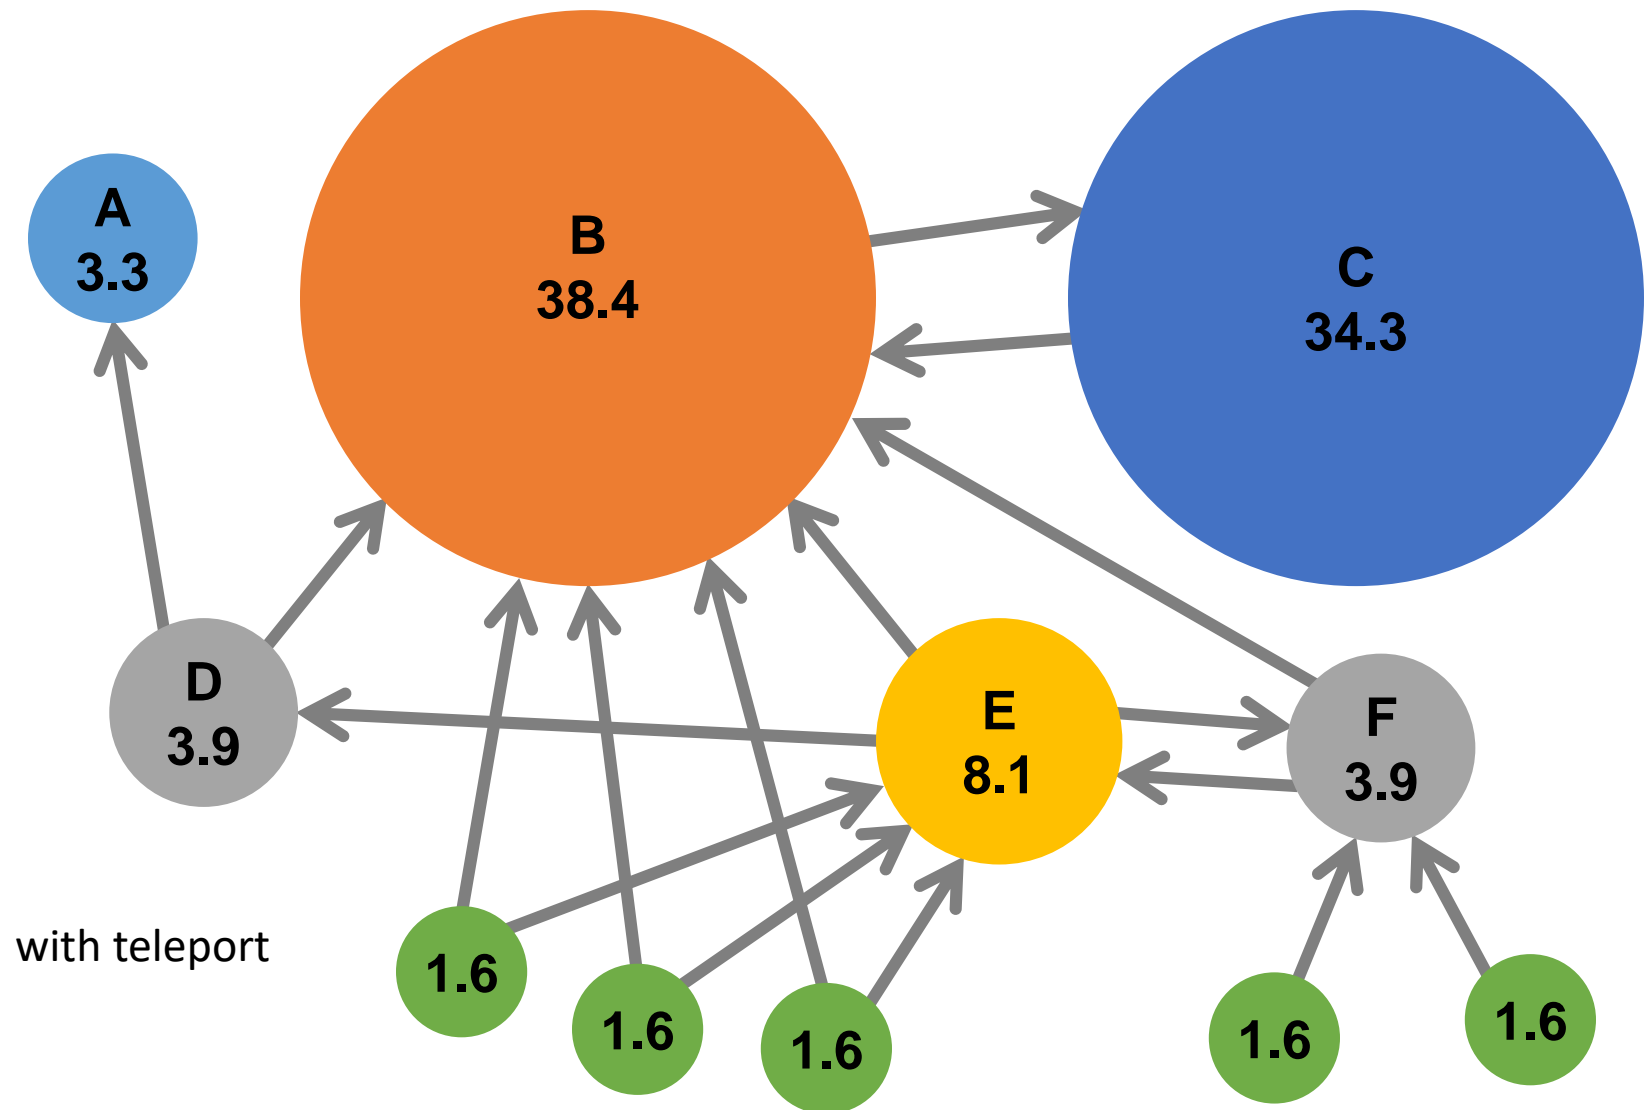
- "Voting" principle
 - each page votes for the importance of the pages it points to
 - a link's vote is proportional to the importance of its source page
 - If page j with importance r_j has n out-links, each link gets $\frac{r_j}{n}$ votes
 - Page j 's own importance is the sum of the votes on its in-links

- Rank of page j : $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
 - d_i ... out-degree of node i

$$r_j = r_i/3 + r_k/4$$



Example: PageRank Scores




Computation via Solving Equations

- Rank of page j : $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$
 - d_i ... out-degree of node i

- Example:
 - 3 equations, 3 unknowns, no constants
 - No unique solution
 - All solutions equivalent modulo a scale factor
 - Additional constraint forces uniqueness: $\sum_i r_i = 1$
 - Solution: $r_y = \frac{2}{5}$, $r_a = \frac{2}{5}$, $r_m = \frac{1}{5}$

- Gaussian elimination method works for small examples but we need a better method for large web-size graphs



Equations:

$$r_y = r_y/2 + r_a/2$$


$$r_a = r_y/2 + r_m$$

$$r_m = r_a/2$$

PageRank: Matrix Formulation

- Stochastic adjacency matrix M
 - If $i \rightarrow j$, then $M_{ji} = \frac{1}{d_i}$ else $M_{ji} = 0$
 - M is a column stochastic matrix
 - Columns sum to 1
- Rank vector r
 - r_i is the importance score of page i
 - $\sum_i r_i = 1$
- Equations $r_j = \sum_{i \rightarrow j} \frac{r_i}{d_i}$ can be written as:

$$r = M \cdot r$$



$$M = \begin{matrix} & \begin{matrix} y & a & m \end{matrix} \\ \begin{matrix} y \\ a \\ m \end{matrix} & \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \end{matrix}$$

$$r = M \cdot r$$

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix}$$

Computation via Eigenvector

- Equations can be written as: $\mathbf{r} = \mathbf{M} \cdot \mathbf{r}$

- The rank vector \mathbf{r} is an eigenvector of the stochastic matrix \mathbf{M}
 - eigenvector with corresponding eigenvalue 1
 - Math background: largest eigenvalue of \mathbf{M} is 1 since \mathbf{M} is column stochastic (with non-negative entries)
 - We know \mathbf{r} is unit length and each column of \mathbf{M} sums to one, so $\mathbf{M}\mathbf{r} \leq \mathbf{1}$

- Finding \mathbf{r} = finding eigenvector of \mathbf{M} corresponding to the largest eigenvalue
 - you know how to do this efficiently (see slides on power iteration)

Notes on Computation

- Power iteration: iteratively compute $\mathbf{r} \leftarrow \frac{\mathbf{M} \cdot \mathbf{r}}{\|\mathbf{M} \cdot \mathbf{r}\|}$ until convergence
 - required for PageRank: $\sum_i r_i = 1$
- Let $\mathbf{y} = \mathbf{M} \cdot \mathbf{x}$ with $\sum_i x_i = 1$.
 Since \mathbf{M} is column stochastic, it holds $\sum_i y_i = 1$
- No need for normalization!
- Start with random (normalized) vector \mathbf{r} , and iterate $\mathbf{r} \leftarrow \mathbf{M} \cdot \mathbf{r}$
- Important: Matrix \mathbf{M} is sparse!
 - we only need to consider the (ingoing) neighbors of each node
- Iteratively compute $r_j \leftarrow \sum_{i \rightarrow j} \frac{r_i}{d_i}$ until convergence
 - first compute the updated value for each r_j , then assign them at once


$$\mathbf{r} = \mathbf{M} \cdot \mathbf{r}$$

$$\begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} r_y \\ r_a \\ r_m \end{bmatrix}$$

Random Walk Interpretation

- Consider a random web surfer that moves between the web pages
 - At time t , the web surfer is in a random webpage i
 - At time $t + 1$, the surfer follows an out-link from i uniformly at random

- The surfer's path (denoted by X_1, X_2, X_3, \dots) forms a Markov chain
 - Web pages are the states of the Markov chain
 - The surfer starts from a random webpage: $\Pr(X_1 = i) = \pi_i$
 - Transition probabilities: $\Pr(X_{t+1} = j | X_t = i) = M_{ji}$
 - Note: the transition probability matrix of the Markov chain is $\mathbf{B} = \mathbf{M}^T$



Random Walk Interpretation

- Under some “technical conditions”, we have that

$$\text{rank score of page } i = r_i = \lim_{t \rightarrow \infty} \Pr(X_t = i)$$

$$\text{or in vector form: } \mathbf{r} = \lim_{t \rightarrow \infty} \boldsymbol{\pi}(t)$$

$$\boldsymbol{\pi}(t) = \boldsymbol{\pi} \mathbf{B}^{(t-1)} \Rightarrow \text{limit of the sequence } \boldsymbol{\pi} \mathbf{B}, (\boldsymbol{\pi} \mathbf{B}) \mathbf{B}, ((\boldsymbol{\pi} \mathbf{B}) \mathbf{B}) \mathbf{B}, \dots \text{ equals to } \mathbf{r}$$

remember

$$\Pr(X_t = i) \stackrel{\text{def}}{=} \pi_i(t)$$

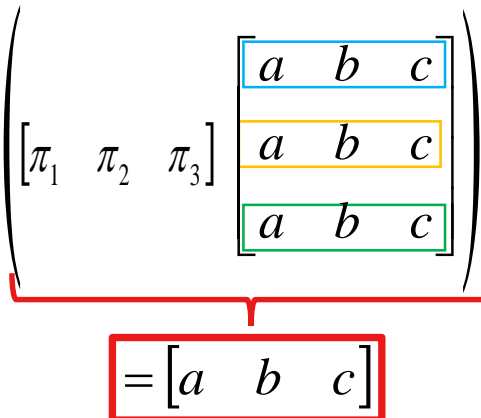
$$\boldsymbol{\pi}(t) = \boldsymbol{\pi} \mathbf{B}^{(t-1)}$$

$$\Pr(X_t = j | X_1 = i) = [\mathbf{B}^{(t-1)}]_{ij}$$

Random Walk Interpretation

- What happens if we do infinitely many steps?
 - $\lim_{t \rightarrow \infty} \boldsymbol{\pi}(t)$ is called the limiting distribution (if it exists)
- Assume that when $t \rightarrow \infty$, \mathbf{B}^t converges to a matrix whose rows are the same.
 - In this case: one row of $\lim_{t \rightarrow \infty} \mathbf{B}^t$ specifies the limiting distribution.
 - And: probability of reaching a node does not depend on start point.

$$\lim_{t \rightarrow \infty} \mathbf{B}^{(t-1)} = \begin{bmatrix} \boxed{a \quad b \quad c} \\ \boxed{a \quad b \quad c} \\ \boxed{a \quad b \quad c} \end{bmatrix} \Rightarrow \lim_{t \rightarrow \infty} \boldsymbol{\pi}(t) = \lim_{t \rightarrow \infty} \boldsymbol{\pi} \mathbf{B}^{(t-1)} = \left(\begin{bmatrix} \pi_1 & \pi_2 & \pi_3 \end{bmatrix} \begin{bmatrix} \boxed{a \quad b \quad c} \\ \boxed{a \quad b \quad c} \\ \boxed{a \quad b \quad c} \end{bmatrix} \right)$$



$$= \boxed{a \quad b \quad c}$$

Random Walk Interpretation

- Stationary distribution: the vector $\pi(\infty)$ is called stationary distribution if the following equality holds


$$\pi(\infty) = \pi(\infty)B$$

- By definition, $\pi(\infty)$ (if exists) is equal to (transpose of) the rank vector r .
- $\pi(\infty)$ can be computed by
 - getting the eigenvector of M associated with the unit eigenvalue
 - normalizing it to one.
- Under the “technical conditions”, a Markov chain has a limiting distribution which is equal to its unique stationary distribution.

Existence and Uniqueness

- What are the “technical conditions”?
 - Being **Irreducible** and **Aperiodic**
- **Irreducible**: it is possible to get to any state from any state
- **Aperiodic**: a state i is aperiodic if there exists n such that for all $n' \geq n$:

$$\Pr(X_{n'} = i | X_1 = i) > 0$$
 - A Markov chain is aperiodic if every state is aperiodic
 - An irreducible Markov chain only needs one aperiodic state to imply all states are aperiodic






PageRank: Problems

- Some pages are dead ends (have no out-links)
 - Random walk has “nowhere” to go to
 - Such pages cause importance to “leak out”


- Spider traps: (all out-links are within the group)
 - Random walk gets “stuck” in a trap
 - And eventually spider traps absorb all importance

- Periodic states:
 - If we start at the state, we will return to the state in fixed periods.



Solution: Random Teleports

- At each step, random surfer has **two options**:
 - With probability β , follow a link at random
 - With probability $1 - \beta$, jump to some random page



- PageRank equation [Brin-Page, 98]

$$r_j = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$$

$$// = \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + \sum_i (1 - \beta) \frac{r_i}{N}$$

- In matrix notation: $A = \beta M + (1 - \beta) \begin{bmatrix} 1/N \\ 1/N \\ \vdots \\ 1/N \end{bmatrix}_{N \times N}$

$[1/N]_{N \times N}$ is a
 N by N matrix
 where all entries
 are $1/N$

- final solution: $\mathbf{r} = \mathbf{A} \cdot \mathbf{r}$

This formulation assumes that \mathbf{M} has no dead ends. We can either preprocess matrix \mathbf{M} to remove all dead ends or explicitly follow random teleport links with probability 1.0 from dead-ends.

Illustration: Random Teleports ($\beta = 0.8$)





Illustration: Random Teleports ($\beta = 0.8$)



Notes on Computation

- Attention: **The matrix \mathbf{A} is dense!**
 - N^2 non-zero entries
 - you should never compute \mathbf{r} in such a way

- Consider the teleport by adding constant penalty to each term
 - iterate $r_j \leftarrow \sum_{i \rightarrow j} \beta \frac{r_i}{d_i} + (1 - \beta) \frac{1}{N}$ until convergence
 - only neighbors need to be considered

- To maintain sparsity in matrix form multiply by $\beta \mathbf{M}$ then add a vector
 - $\mathbf{r} = \beta \mathbf{M} \mathbf{r} + (1 - \beta) \left[\frac{1}{N} \right]_N$

- **Vertex-oriented computation**
 - each vertex performs local computations

Systems/Frameworks for Graph Processing

- Specialized systems for such kind of graph processing
 - *GraphLab* (Dato, Turi)
 - *Giraph* (open source counterpart to Google's Pregel)
 - *GraphX*: Library for graph processing on top of Spark
- **Crucial aspect: vertex-oriented programming**
 - each vertex performs local computations
 - GAS principle — **gather, apply, scatter**: each vertex (a) gathers information from adjacent vertices/edges (b) applies transformation, (c) scatters information to adjacent vertices
 - for PageRank only steps a + b required
- Similar concepts become also more frequent in Deep Learning Frameworks due to popularity of Graph Neural Networks

Some Problems with Page Rank

- Measures generic popularity of a page
 - Biased against topic-specific authorities
 - Solution: Topic-Sensitive PageRank
- Susceptible to Link spam
 - Artificial link topographies created in order to boost PageRank
 - Solution: TrustRank
- Uses a single measure of importance
 - Other models of importance
 - Solution: Hubs-and-Authorities

Topic-Sensitive PageRank

- Instead of **generic popularity**, can we measure popularity **within a topic**?
 - Goal: Evaluate Web pages not just according to their popularity, but by how close they are to a particular topic, e.g. “sports” or “history”
 - Allows search queries to be answered based on **interests of the user**

- Core idea: **Bias the random walk**
 - When walker teleports, pick a page from a set S
 - **Standard PageRank**: S = all pages
 - any page with equal probability
 - **Topic-Sensitive PageRank**: S = set of “relevant” pages
 - E.g., Open Directory (DMOZ) pages for a given topic/query
 - For each teleport set S , we get a different vector \mathbf{r}_S

Generalizing Topic-Sensitive PageRank

- As a matrix equation topic-sensitive PageRank takes the following form


$$r = \beta M r + (1 - \beta) \pi \quad \text{where } \pi_i = \begin{cases} \frac{1}{|S|} & \text{if } i \in S \\ 0 & \text{otherwise} \end{cases}$$

- We can generalize this further to arbitrary teleport vectors π

$$r = \beta M r + (1 - \beta) \pi \quad \text{where } \sum_i \pi_i = 1$$

- The exact solution is $r = (1 - \beta)(I - \beta M)^{-1} \pi$
 - Runtime scales worse than $O(N^2)$
 - Use the iterative approximate algorithm in practice
 - Multiply by $\beta \cdot M$, then add restart vector $(1 - \beta)\pi$, repeat, ...
 - Maintains sparsity

Example: Topic-Sensitive PageRank



Suppose $S = \{1\}$, $\beta = 0.8$

Node	Iteration				
	0	1	2	...	stable
1	0.25	0.4	0.28		0.294
2	0.25	0.1	0.16		0.118
3	0.25	0.3	0.32		0.327
4	0.25	0.2	0.24		0.261

$S = \{1\}$, $\beta = 0.90$:
 $r = [0.17, 0.07, 0.40, 0.36]$
 $S = \{1\}$, $\beta = 0.8$:
 $r = [0.29, 0.11, 0.32, 0.26]$
 $S = \{1\}$, $\beta = 0.70$:
 $r = [0.39, 0.14, 0.27, 0.19]$

$S = \{1,2,3,4\}$, $\beta = 0.8$:
 $r = [0.13, 0.10, 0.39, 0.36]$
 $S = \{1,2,3\}$, $\beta = 0.8$:
 $r = [0.17, 0.13, 0.38, 0.30]$
 $S = \{1,2\}$, $\beta = 0.8$:
 $r = [0.26, 0.20, 0.29, 0.23]$
 $S = \{1\}$, $\beta = 0.8$:
 $r = [0.29, 0.11, 0.32, 0.26]$

Discovering the Topic Set S

- **Create different PageRanks for different topics**
 - The 16 DMOZ top-level categories:
 - arts, business, sports,...

- **Which topic ranking to use?**
 - User can pick from a menu
 - Classify query into a topic
 - Can use the **context** of the query
 - E.g., query is launched from a web page talking about a known topic
 - History of queries e.g., “basketball” followed by “Jordan”
 - User context, e.g., user’s bookmarks, ...

PageRank: Variants (I)

- **“Normal” PageRank:**

- Teleports uniformly at random to any node
- All nodes have the same teleport probability of surfer landing there:

$$\pi = (0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1 \quad 0.1)^T$$

- **Topic-Sensitive PageRank:**

- Teleports to a topic specific set of pages
- Nodes can have different probabilities of surfer landing there:

$$\pi = (0.1 \quad 0 \quad 0 \quad 0.2 \quad 0 \quad 0.5 \quad 0 \quad 0 \quad 0 \quad 0.2)^T$$

- **Personalized PageRank (Random Walk with Restarts):**

- Teleport is always to the same node:

$$\pi = (0 \quad 0 \quad 0 \quad 1 \quad 0 \quad 0 \quad 0 \quad 0 \quad 0 \quad 0)^T$$

PageRank: Variants

- Spam is common in the web
 - Spammer's goal: Maximize the PageRank of target page t
 - Technique:
 - Get as many links from accessible pages as possible to target page t
 - Construct “link farm” to get PageRank multiplier effect


- Combating link spam via TrustRank
 - **Topic-sensitive PageRank with a teleport set of trusted pages**
 - Example: .edu domains, similar domains for non-US schools

Summary

- Core idea: Ranking of the nodes based on the link structure
- PageRank scores nodes depending on their incoming links
- With a teleport set we can rank nodes based on arbitrary factors, for example
 - Topic
 - Trust
 - Node identity
- Computing PageRank requires sparse matrix products for even moderately sized graphs

Questions

- Consider a directed cycle of length 3 as a Markov chain disregarding edge weights



- Is it irreducible? Is it aperiodic?
- How does the introduction of random teleports change the above 3-cycle?
- How can you make it aperiodic by inserting just a single edge?

Machine Learning for Graphs and Sequential Data

Graphs – Ranking

Lecturer: Prof. Dr. Stephan Günnemann

www.daml.in.tum.de

Summer Term 2020

Data Analytics and
Machine Learning 