

# Machine Learning for Graphs and Sequential Data

## *Graphs – Semi-Supervised Learning*

Lecturer: Prof. Dr. Stephan Günnemann

[www.daml.in.tum.de](http://www.daml.in.tum.de)

---

Summer Term 2020

Data Analytics and  
Machine Learning 

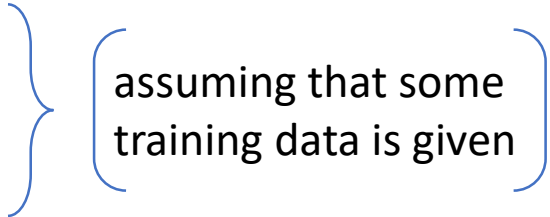
# Roadmap

---

- **Chapter: Graphs**

1. Graphs & Networks
2. Generative Models
3. Clustering
4. Node Embeddings
5. Ranking
- 6. Semi-Supervised Learning**
  - **Label Propagation**
  - Graph Neural Networks
7. Limitations of GNNs


# Types of Machine Learning Problems on Graphs

- So far we have discussed unsupervised learning problems on graphs
  - generative modeling
  - clustering / community detection
  - node embeddings
  - ranking
- What about supervised learning tasks, such as
  - classifying role of a protein in a PPI network
  - detecting fraudsters in an e-commerce system
  - predicting user's preferences in a social network

assuming that some training data is given
- More generally, how do we label/categorize/classify instances in a graph?

# Collective Classification



- Consider the following problem
  - Graph represents a social network, nodes = users, edges = friendship
  - Labels are known for some **labeled nodes**
  - Goal is to classify the **unlabeled nodes**



- Standard assumption: **Homophily** (a.k.a. **assortativity**)
  - *"birds of a feather flock together"*
  - *a.k.a. smoothness assumptions*
  - *that is, if nodes are connected by an edge, they are likely to have same labels*

# Label Propagation

- Consider the binary case (two classes)
- Formal definition of the problem
  - Nodes  $V = S \cup U$ 
    - Labeled instances  $S$  (seeds) and unlabeled instances  $U$
  - Symmetric weighted adjacency matrix  $\mathbf{W} \in \mathbb{R}^{|V| \times |V|}$ 
    - $w_{ij} \geq 0$  denotes similarity of nodes  $i$  and  $j$
  - $\hat{y}_i \in \{0,1\}$  for  $i \in S$  // **given** class labels for the nodes in  $S$
  - $y_i \in \{0,1\}$  for  $i \in V$  // class labels we want to **predict** for each node
- Idea: **Energy minimization**  $E(\mathbf{y}) = \frac{1}{2} \sum_{ij} w_{ij} (y_i - y_j)^2$ 
  - **smoothness**: adjacent nodes should have same class label



# Energy Minimization as MAP Inference

- Goal: find the labeling  $\mathbf{y} \in \{0,1\}^N$  that minimizes the energy

$$\min_{\mathbf{y}} E(\mathbf{y}) = \min_{\mathbf{y}} \frac{1}{2} \sum_{ij} w_{ij} (y_i - y_j)^2$$

- This is equivalent to the optimization problem

$$\arg \min_{\mathbf{y} \in \{0,1\}^N} E(\mathbf{y}) = \arg \max_{\mathbf{y} \in \{0,1\}^N} \exp(-E(\mathbf{y}))$$

- Consider the following approximation to the posterior

$$p(\mathbf{y}_U | \mathbf{W}, \mathbf{y}_L) = \frac{1}{Z} \exp(-E(\mathbf{y}))$$

where  $Z = \sum_{\mathbf{y}} \exp(-E(\mathbf{y}))$  is the normalizing constant.

- Energy minimization is equivalent to **MAP inference** in this model!
- Such approaches are called “Energy-Based Learning”.
- In general: any nonnegative energy function  $E(\mathbf{y})$  corresponds to a probability distribution over  $\mathbf{y}$  (if  $\text{dom}(\mathbf{y})$  is finite, restrictions apply in other cases).

# Label Propagation

- Two aspects: **Smoothness** + **Matching the seed labels**


$$\min_{\mathbf{y} \in \{0,1\}^{|V|}} \frac{1}{2} \sum_{i,j} w_{ij} (y_i - y_j)^2 \text{ subject to } y_i = \hat{y}_i \text{ for all } i \in S$$

- Constrained integer optimization problem

- We know how to rewrite the above problem!

$$\min_{\mathbf{y} \in \{0,1\}^{|V|}} \mathbf{y}^T \mathbf{L} \mathbf{y} \text{ subject to } y_i = \hat{y}_i \text{ for all } i \in S$$

- $\mathbf{L} = \mathbf{D} - \mathbf{W}$  is the graph Laplacian
- And we know how to make it more tractable
  - Drop the integer constraint:  $y_i$  ( $i \in U$ ) can be any real value



# Label Propagation: Solution

- Task:  $\min_{\mathbf{y} \in \mathbb{R}^{|V|}} \mathbf{y}^T \mathbf{L} \mathbf{y}$  subject to  $y_i = \hat{y}_i$  for all  $i \in S$

- Solution:

- w.l.o.g. assume the Laplacian matrix is partitioned into blocks for labeled and unlabeled nodes

$$\mathbf{L} = \begin{bmatrix} \mathbf{L}_{SS} & \mathbf{L}_{SU} \\ \mathbf{L}_{US} & \mathbf{L}_{UU} \end{bmatrix}$$


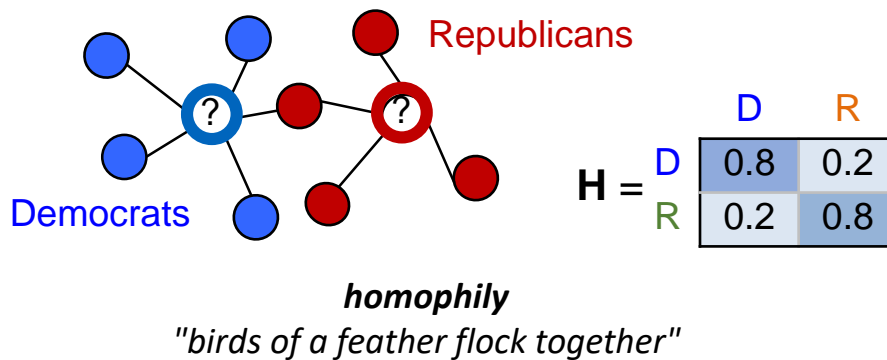
- Accordingly let  $\mathbf{y} = \begin{bmatrix} \mathbf{y}_S \\ \mathbf{y}_U \end{bmatrix} = \begin{bmatrix} \hat{\mathbf{y}}_S \\ \mathbf{y}_U \end{bmatrix}$  the vector of labels to be learned
- Then:  $\mathbf{y}_U = -\mathbf{L}_{UU}^{-1} \cdot \mathbf{L}_{US} \cdot \hat{\mathbf{y}}_S$

[Zhu2002]



# Label Propagation: Generalization

- What if we have  $K$  labels?
  - Use one-hot notation  $y_{ik} = \begin{cases} 1 & \text{if node } i \text{ is of class } k \\ 0 & \text{else} \end{cases}$
  - Energy function  $E(Y) = \sum_{i,j} w_{ij} (\mathbf{y}_i - \mathbf{y}_j)^T (\mathbf{y}_i - \mathbf{y}_j)$
- Other types of **network effects** – encode with a compatibility matrix  $\mathbf{H}$



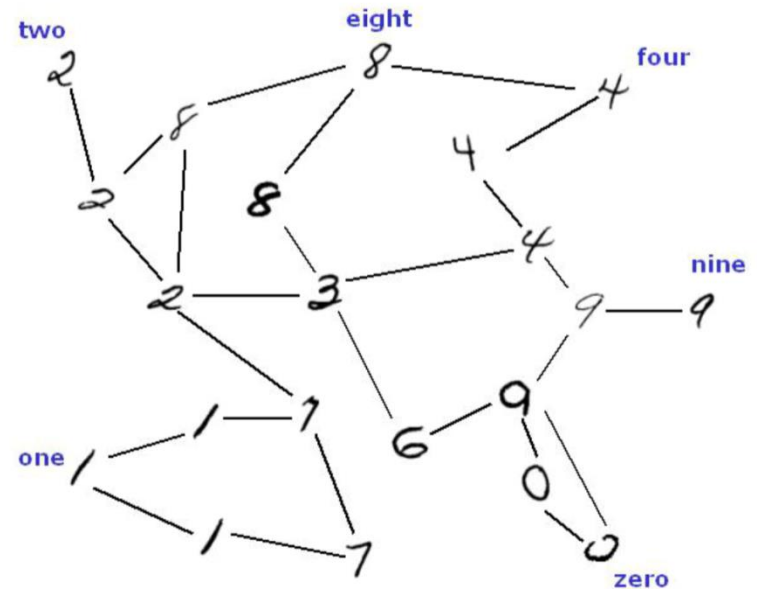
- Energy function  $E(Y) = \sum_{i,j} w_{ij} (\mathbf{y}_i - \mathbf{y}_j)^T \mathbf{H} (\mathbf{y}_i - \mathbf{y}_j)$

# Label Propagation vs. SBM

- At first glance both models seem very similar
  - labels  $\mathbf{y}_i$  look a lot like community affiliations  $\mathbf{z}_i$
  - compatibility matrix  $\mathbf{H}$  from LP looks like  $\boldsymbol{\eta}$  from SBM
  
- Is LP equivalent to inference in SBM with some  $\mathbf{z}_i$ s observed?
  - Label propagation is a discriminative model that only models the conditional distribution of labels given the similarity graph  $p(\mathbf{Y}|\mathbf{W})$
  - on the other hand, SBM is a generative model that models  $\Pr(\mathbf{A}|\mathbf{Z})$  and  $\Pr(\mathbf{Z})$ 
    - we can use SBM to generate new graphs – not the case for LP!
  - for SBM we get the posterior  $\Pr(\mathbf{Z}|\mathbf{A}) = \frac{\Pr(\mathbf{A}|\mathbf{Z}) \Pr(\mathbf{Z})}{\Pr(\mathbf{A})}$  using Bayes' formula
  
- SBM and LP solve different problems
  - SBM: estimate what parameters generated a given graph  $\mathbf{A}$  (unsupervised)
  - LP: predict labels of the nodes in  $U$  given observed labels and  $\mathbf{W}$  (supervised)

# Non-Graph Data

- What if we only have vector data (no graph is available)?
- Simply construct a graph connecting similar data points
  - graph construction: see section on spectral clustering (e.g. k-NN graph)
- Apply label propagation just like before.
- Features are used to construct the graph, but then only the graph is used for classification.



# Transductive Learning

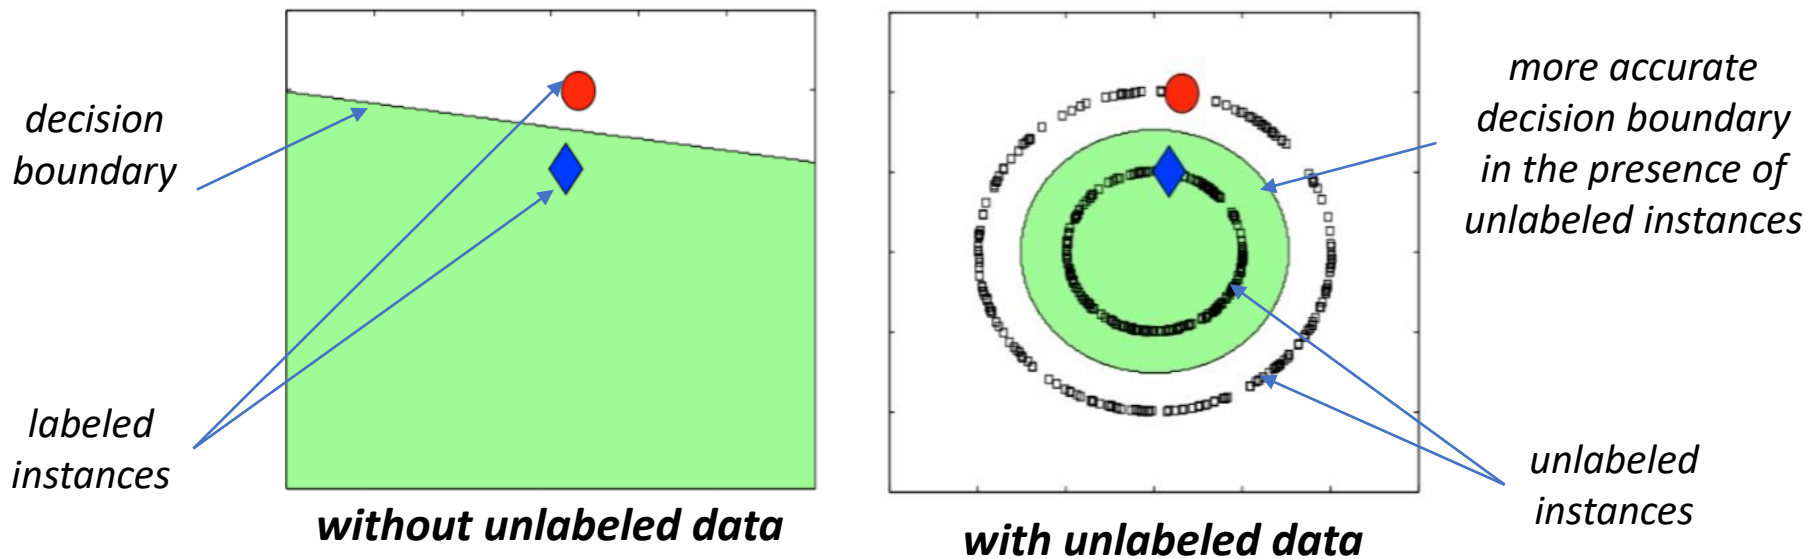
- Label Propagation is a special case of so-called **transductive learning**.
  - Given
    - (i) a set of labeled training instances  $T = \{(x_i, y_i)_{i=1 \dots N}\} \subset \mathcal{X} \times \mathcal{Y}$
    - (ii) a set of unlabeled test instances  $U = \{(x_i)_{i=1 \dots M}\} \subset \mathcal{X}$
    - [+ potentially some other knowledge (graph structure  $\mathbf{W}$ , affinity matrix  $\mathbf{H}$ )]
  - Goal
    - predict labels **only** for the unlabeled instances  $U$  (i.e. learn  $f: U \rightarrow \mathcal{Y}$ )
  - *“When trying to solve some problem, one should not solve a more difficult problem as an intermediate step”* – Vapnik’s principle
- “Traditional” supervised learning (e.g. NN, SVM) is **inductive learning**:
  - Given
    - (i) a set of labeled training instances  $T = \{(x_i, y_i)_{i=1 \dots N}\} \subset \mathcal{X} \times \mathcal{Y}$
    - [+ potentially some other knowledge]
  - Goal
    - learn a prediction function (mapping)  $f: \mathcal{X} \rightarrow \mathcal{Y}$  (that can be applied to any  $x_{new} \in \mathcal{X}$ )

# Transduction vs. Semi-Supervised Learning

- LP in graphs is often referred to as "graph-based semi-supervised learning"
  - not a complete misnomer, but a more specific term would be: graph-based transductive learning
- **Semi-supervised learning (SSL)** is a more generic principle.
- Standard definition:
  - Given: labeled data  $T = \{(x_i, y_i)_{i=1 \dots N}\}$  and unlabeled data  $U = \{(x_i)_{i=1 \dots M}\}$ .
  - Main idea: Use **both**  $T$  **and**  $U$  to learn a mapping  $f$ .  
This can be either inductive ( $f: \mathcal{X} \rightarrow \mathcal{Y}$ ) or transductive ( $f: U \rightarrow \mathcal{Y}$ ).
- Transductive learning is almost always semi-supervised:
  - We are given  $T$  and  $U$ . The goal is to predict labels only for  $U$ .
  - Of course we will use  $U$  to do this!  $\Rightarrow$  Semi-supervised learning

# Why Semi-Supervised Learning Works?

- How can unlabeled data be helpful?
  - Unlabeled data helps us to better model the data distribution

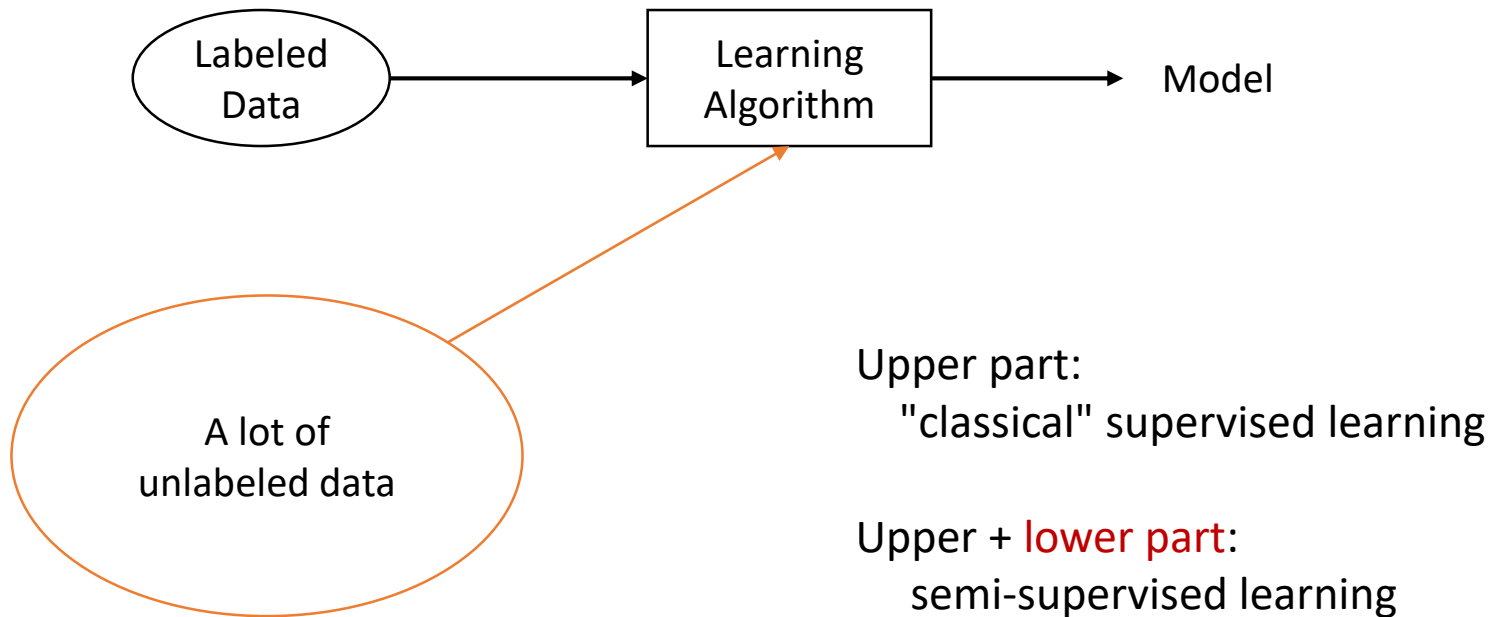


- Caveat:
  - We need to make assumptions about the data/label distribution (e.g. manifold / smoothness / cluster / low-density separation assumptions)
  - If the assumptions are wrong, SSL may perform even worse than simple SL!

Example from [Belkin et al., JMLR 2006]

# Semi-supervised Learning: Motivation

- Why semi-supervised learning?
- Large amounts of unlabeled data, small amounts of labeled data
- Labeling/annotating data is expensive



# Roadmap

---

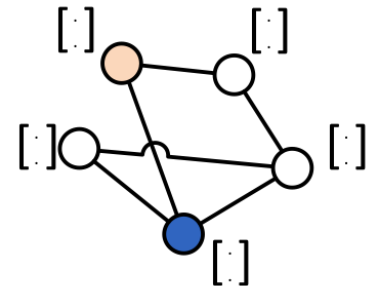
- **Chapter: Graphs**

1. Graphs & Networks
2. Generative Models
3. Clustering
4. Node Embeddings
5. Ranking
- 6. Semi-Supervised Learning**
  - Label Propagation
  - **Graph Neural Networks**
7. Limitations of GNNs




# (Semi-supervised) Deep Learning on Graphs

- Neural Networks have achieved outstanding performance for many data types
  - However: usually restricted to simple grid-like (images) or sequential data (text)
- How about neural network for graphs?
- Additional motivation: In a lot of real world graphs the nodes have attributes
  - In citation networks nodes are papers, the text gives rise to attributes, and edges are citations; in protein-protein interaction networks the properties of the proteins can be considered as attributes
  - Label propagation, however, considers only the network structure
  - How can we perform semi-supervised learning on graphs taking both the network structure and the attributes into account?
- Idea: Differentiable message passing
  - a.k.a. neural message passing, graph convolutions, (or more general: graph neural networks)



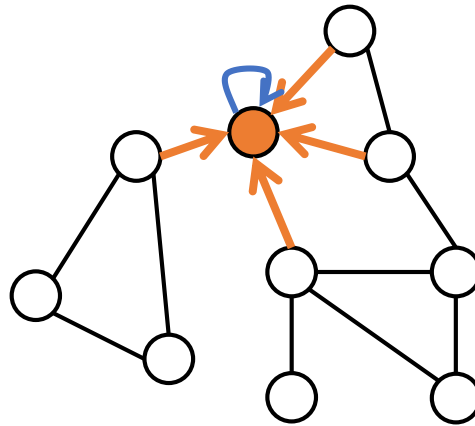
# Differentiable Message Passing Framework (I)



- Given:
  - A graph  $G = (V, E)$  with a set of nodes  $V$ ,  $E$  set of edges,  $E_{vu}$  edge weight
  - Each node  $v$  has features  $x_v \in \mathbb{R}^d$
  - A subset of labeled nodes  $S \subseteq V$ , where  $y_v$  denotes the label of node  $v \in S$
  
- Let  $h_v^{(k-1)}$  be the hidden representation of node  $v$  at previous  $k - 1$  layer
  
- For each node do:
  1. Gather messages from all neighbors  $m_v^{(k)} = \sum_{u \in N(v)} M(h_v^{(k-1)}, h_u^{(k-1)}, E_{vu})$
  2. Update the hidden representation  $h_v^{(k)} = U(h_v^{(k-1)}, m_v^{(k)})$
  
- $M$  and  $U$  are any differentiable functions, e.g. neural networks

# Differentiable Message Passing Framework (II)

- For each node do:
  1. Gather messages from all neighbors  $m_v^{(k)} = \sum_{u \in N(v)} M(h_v^{(k-1)}, h_u^{(k-1)}, E_{vu})$
  2. Update the hidden representation  $h_v^{(k)} = U(h_v^{(k-1)}, m_v^{(k)})$
- Example: Calculating the update for the node in orange



1. Gather message from all neighbors
2. Update hidden representation

# Example Instantiation of the Framework

- Let  $h_v^{(0)} = x_v$ 
  - At the first layer the representations are the node features
- Set the message aggregation function to be the average over the neighbors

$$m_v^{(k)} = \sum_{u \in N(v)} \frac{1}{d_v} (W^{(k)} h_u^{(k-1)} + b^{(k)})$$

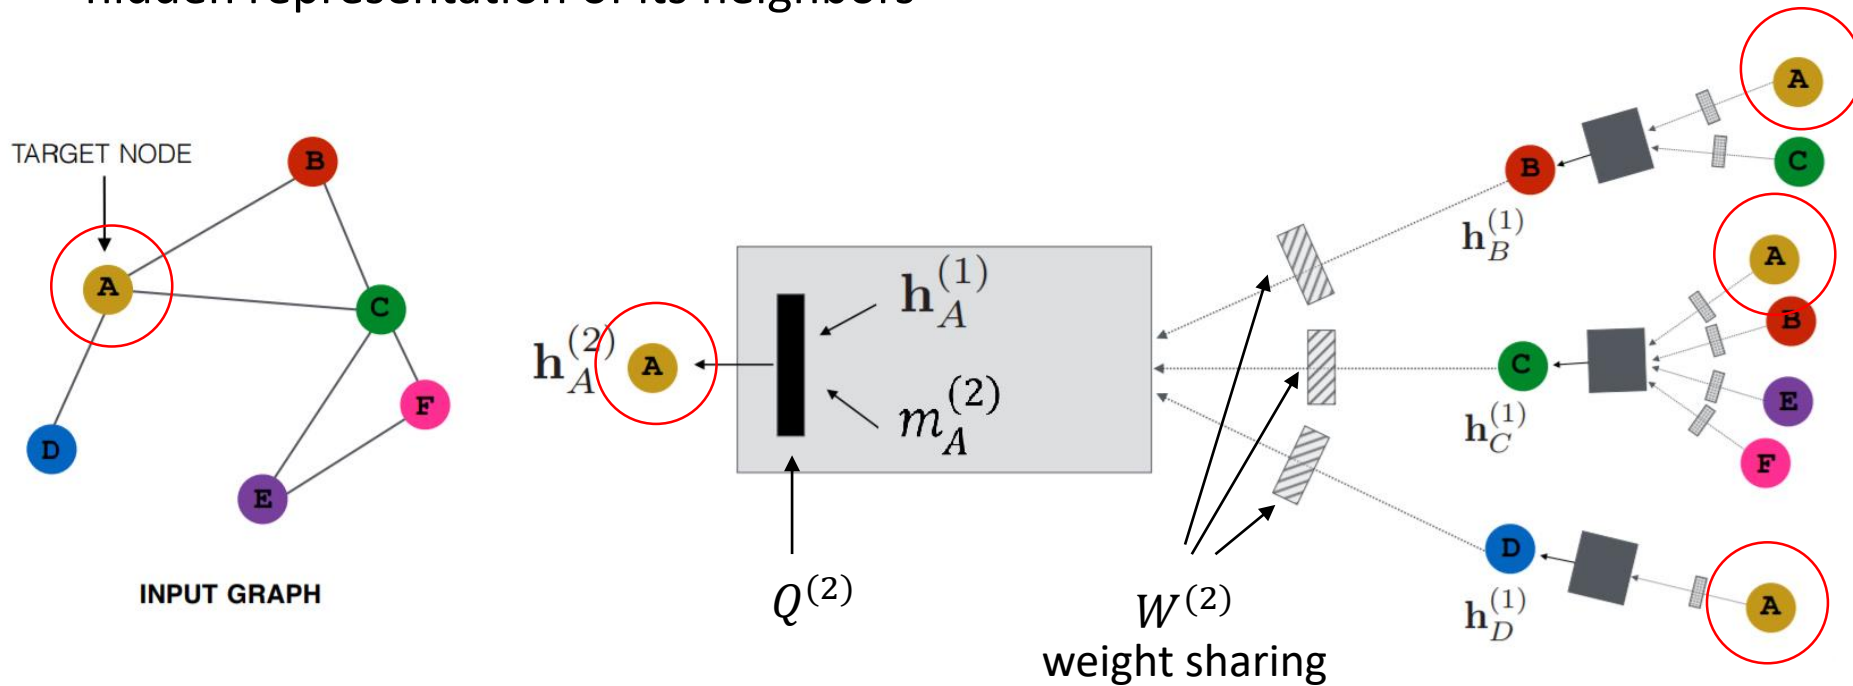
- Calculate hidden representations as simple NNs with a non-linearity

$$h_v^{(k)} = \text{relu}(Q^{(k)} h_v^{(k-1)} + p^{(k)} + m_v^{(k)})$$

- Here  $W^{(k)}, b^{(k)}, Q^{(k)}, p^{(k)}$  are the **trainable parameters** of the  $k$ -th layer

# Recursive View of Differentiable Message Passing

- The hidden representation of each node is recursively defined in terms of the hidden representation of its neighbors



# How do Neighbors Influence a Given Node?

- Observation 1:  $K$  hidden layers equal to  $K$  steps of message passing
  - At step 1, node  $v$  aggregates information from its 1-hop neighbors
  - At step 2, node  $v$  (implicitly) aggregates information from the neighbors of its 1-hop neighbors, i.e. from its 2-hop neighbors
  - ....
  - At step  $K$ , node  $v$  (implicitly) aggregates information from its  $K$ -hop neighbors
- Observation 2:  $K$  hidden layers mean that the representation  $h_v^{(K)}$  of node  $v$  is based on information from all nodes in its  $K$ -hop neighborhood

# How to Perform Semi-Supervised Node Classification?

- Consider the representations  $h_v^{(K)}$  at the final layer ( $K$ ) as logits
- And use the softmax function to obtain the probability of node  $v$  to belong to class  $c$

$$p_v = \text{softmax}\left(h_v^{(K)}\right)$$

- Train the network using the standard cross entropy loss between the predicted probabilities and the observed labels
  - Recall:  $y_{vc}$  is the one-hot vector encoding the label for node  $v$
  - Denote with  $p_{vc}$  the probability that node  $v$  to belong to class  $c$
  - Denote with  $\mathcal{C}$  the set of all classes, and  $S$  the set of labeled nodes

$$\min_{\{W^{(k)}, Q^{(k)}, p^{(k)}, b^{(k)}\}_{k=1..K}} - \sum_{v \in S} \sum_{c \in \mathcal{C}} y_{vc} \log p_{vc}$$

# Graph Neural Networks

---

- Exploding interest in recent years:
  - Graph Neural Network (Gori et al., 2005)
  - Spectral Networks and Locally Connected Networks ... (Bruna et al., 2014)
  - Gated Graph Neural Network (Li et al., 2016)
  - Convolutional Neural Networks on Graphs with Fast ... (Defferrard et al., 2017)
  - Neural Message Passing for Quantum Chemistry (Gilmer et al., 2017)
  - Semi-Supervised Classification with Graph Convolutional Nets (Kipf et al., 2017)
  - Graph Attention Networks (Veličković et al., 2018)
  - Predict then Propagate: Graph Neural Networks meet Personalized PageRank (Klicpera et al., 2019)
- Many tasks beyond semi-supervised node classification
  - graph classification, recommendation, 3d-shape-matching, etc.
- Sometimes called geometric deep learning





# Graph Neural Networks

---

- Increasing interest in recent years:
  - Graph Neural Network (Gori et al., 2005)
  - Spectral Networks and Locally Connected Networks ... (Bruna et al., 2014)
  - Gated Graph Neural Network (Li et al., 2016)
  - **Convolutional** Neural Networks on Graphs with Fast ... (Defferrard et al., 2017)
  - Neural Message Passing for Quantum Chemistry (Gilmer et al., 2017)
  - Semi-Supervised Classification with Graph **Convolutional** Nets (Kipf et al., 2017)
  - Graph Attention Networks (Veličković et al., 2018)
  - Predict then Propagate: Graph Neural Networks meet Personalized PageRank (Klicpera et al., 2019)
- Many tasks beyond semi-supervised node classification
  - graph classification, recommendation, 3d-shape-matching, etc.
- Sometimes called geometric deep learning

# From Matrix Convolution ...

- Alternative view of message passing framework is that we are essentially performing a “graph convolution”
- Recap: CNNs

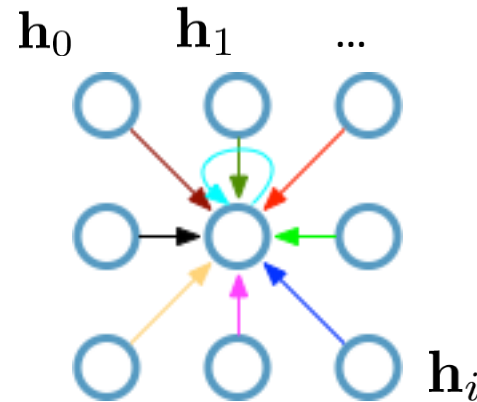



Update for a single pixel:

- Transform messages individually  $\mathbf{W}_i \mathbf{h}_i$
  - Add everything up  $\sum_i \mathbf{W}_i \mathbf{h}_i$
- Images are a special kind of graph: every pixel is a node connected to 8 other nodes (up, down, left, right, etc. pixels)

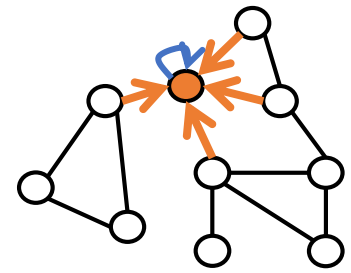
## ... to Graph Convolution

- Alternative view of message passing framework is that we are essentially performing a “graph convolution”
  - Note: Unlike sequences/images, where the convolution operation is clearly defined, there are various versions for graphs. Indeed, one often distinguishes between spatial and spectral approaches.



- Images are a special kind of graph: every pixel is a node connected to 8 other nodes (up, down, left, right, etc. pixels)

# Graph Convolutional Neural Networks



- In the spatial domain, a graph convolution updates each nodes' features by considering the local neighborhood.
  - In effect it is some kind of message passing
  - Example:  $X_i^{(t+1)} = X_i^{(t)} + \sum_{(i,j) \in E} X_j^{(t)}$   
or in matrix notation  $X^{(t+1)} = X^{(t)} + AX^{(t)} = (A + I_n)X^{(t)}$
- Similar to a (learnable) convolutional layer we can formulate a (learnable) graph convolutional layer
 
$$H^{(l+1)} = \sigma \left( \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)} \right)$$

● Non-linearity    ● Message Passing    ● Feature Transformation
- Normalizing the propagation matrix avoids exploding gradients
  - Choose  $\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$  where  $\tilde{A} = A + I_n$  and  $D_{ii} = \sum_j \tilde{A}_{ij}$  is the degree matrix of  $\tilde{A}$

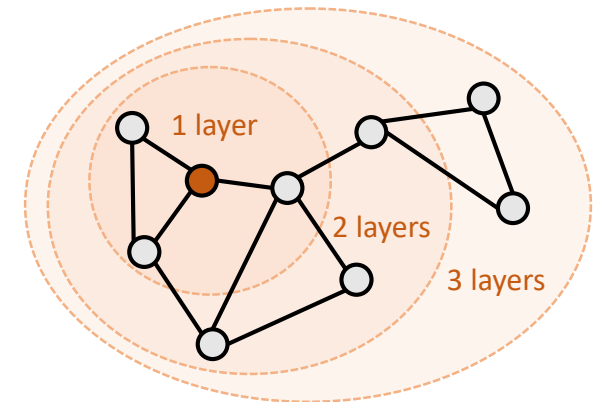
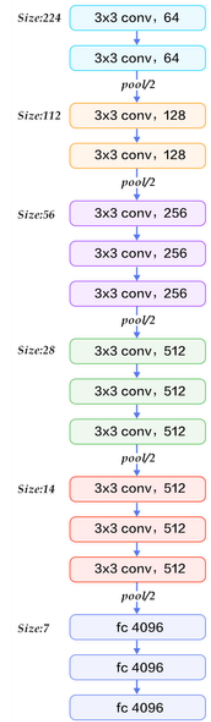
[Kipf2017]

# How Deep are Graph Neural Networks?

- In an MLP or CNN, depth means the number of non-linear transformations of the input
- On graphs we can also consider the furthest distance that a model propagates information as depth

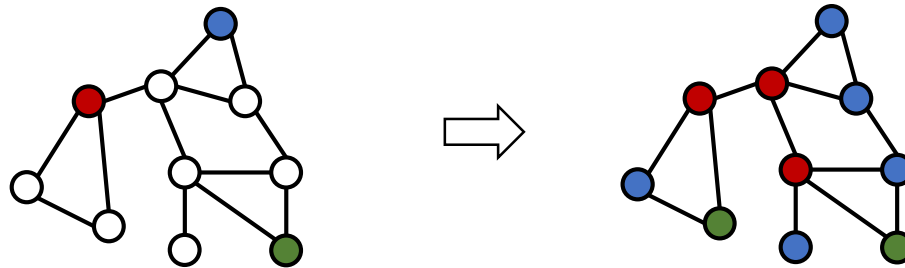
$$H^{(l+1)} = \sigma\left(\tilde{D}^{-\frac{1}{2}}\tilde{A}\tilde{D}^{-\frac{1}{2}}H^{(l)}W^{(l)}\right)$$

- GCNs apply 1 transformation per message passing step
- *Transformation depth* and *propagation depth* are orthogonal in GNNs
  - We could transform the messages with a multi-layer network in each step or propagate without transforming for multiple steps

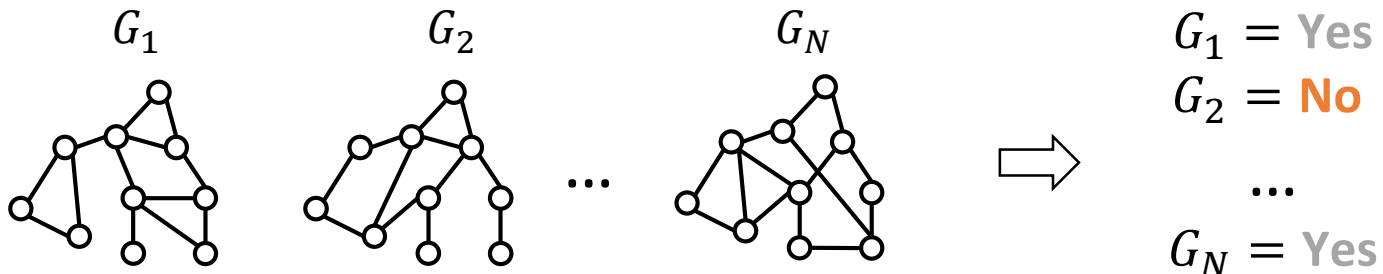


# Single vs. Multi-graph Learning

- So far we had a single graph  $G = (V, E)$  and we learned targets for the nodes
  - For example predict the classes for each node (red, green or blue) in a single large graph



- What if we have multiple graphs as input and the target is for the graph?
  - For example each input is a molecule (i.e. a graph of atoms) and the graph level target is whether it is an effective drug against some disease



# Multi-graph Learning Framework

- Given:
  - A set of graphs  $\mathcal{G} = \{G_i = (V_i, E_i)\}_{i=1..N}$
  - A subset of labeled graphs  $\mathcal{H} \subseteq \mathcal{G}$ , with  $y_{G_i}$  denotes the graph level target of the graph  $G_i \in \mathcal{H}$
- Obtain the hidden representation of the last layer for all nodes in all graphs according to the message passing framework
  - Let  $H_{G_i} = [h_1^{(K)}, h_2^{(K)}, \dots, h_{|V_i|}^{(K)}]$  be a matrix where we stacked the final representations **for all nodes** for graph  $G_i$
- Define an aggregation function  $R(H_{G_i})$  over the node representations of a given graph that produces a representation for the **entire** graph  $h_{G_i} = R(H_{G_i})$
- For classification: consider  $h_{G_i}$  as the logits and train a model using standard cross-entropy loss, i.e.  $\mathcal{L}(y_{G_i}, \text{softmax}(h_{G_i}))$ 
  - For regression, e.g. predicting functional properties of molecules you can use squared loss
- Example agg. function  $R(H_{G_i}) = \frac{1}{|V_i|} \sum_{v \in V_i} h_v^{(K)}$  is the average of the node embedding

# Summary

---

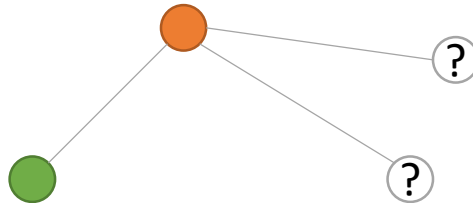
- Semi-supervised learning / graph-based transductive learning
  - Leverage unlabeled data to improve performance of supervised learning
  - Helps if assumptions about the data distribution are correct, e.g. homophily
- Label Propagation spreads labels along the edges of a graph by minimizing the difference between neighbors
  - Usually assumes smoothness but other kinds of network effects can be modeled as well
- Differentiable message passing:
  - Flexible framework to apply the power of deep learning to graphs
  - The nodes aggregate information from their k-hop neighbors
  - Message passing is a generalization of convolution from grids to general graphs
- We distinguish between node-level learning (e.g. node classification) and graph-level learning (e.g. graph classification)



# Questions

---

- Consider the graph below. What is the influence of the green node on the unlabeled nodes in Label Propagation? Why? How about GNNs with  $K = 2$ ?



- Does semi-supervised learning exist outside of learning on graphs?
- What could be alternative aggregation functions beside summation in the Gather step of GNNs?
- Can you apply GNNs to vector data without a specified graph structure?

# Reading Material

---

- [Zhu2002] Zhu, X., & Ghahramani, Z. (2002). Learning from labeled and unlabeled data with label propagation. Center for Automated Learning and Discovery, CMU: Carnegie Mellon University, USA.
- [Kipf2017] Kipf, T. N., & Welling, M. (2017). Semi-Supervised Classification with Graph Convolutional Networks. In ICLR 2017 : International Conference on Learning Representations 2017