

# Machine Learning for Graphs and Sequential Data


## *Graphs – Limitations of GNNs*

Lecturer: Prof. Dr. Stephan Günnemann

[www.daml.in.tum.de](http://www.daml.in.tum.de)

---

Summer Term 2020

Data Analytics and  
Machine Learning 

# Roadmap

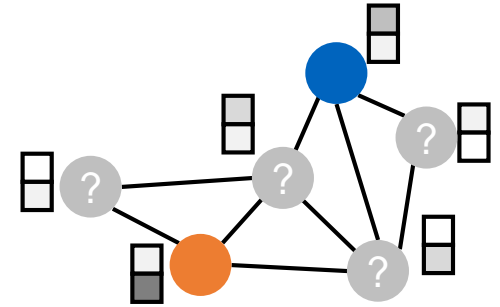
---

- **Chapter: Graphs**

1. Graphs & Networks
2. Generative Models
3. Clustering
4. Node Embeddings
5. Ranking
6. Semi-Supervised Learning
- 7. Limitations of GNNs**
  - Overview
  - **Robustness**

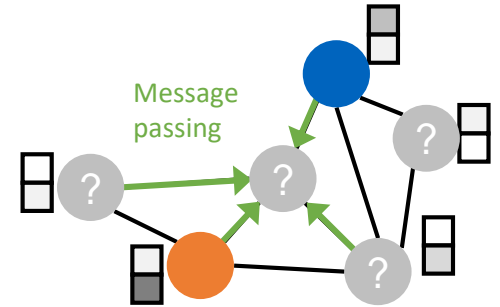
# Adversarial Attacks on GNNs

- Earlier in this course we have seen the problem of (adversarial) robustness of classifiers on “traditional” data, e.g. images.
- In contrast, graph neural networks (GNNs) use **both** the **node’s attributes** as well as their **connections** to make a prediction.
  - Therefore, adversarial attacks can happen through both the **node attributes** as well as the **graph structure**.
  - Structural attacks are indeed quite common in the real world (e.g. adding fake connections in a social network)
- Structure attacks are specifically challenging since they change the flow of messages passed through the GNN



# Adversarial Attacks on GNNs

- Example: two-layer GCN in matrix form:



node attributes

$$\mathbf{Z} \in \mathbb{R}^{N \times C} = f_{\theta}(\mathbf{A}, \mathbf{X}) = \text{softmax}(\hat{\mathbf{A}} \text{ReLU}(\hat{\mathbf{A}} \mathbf{X} \mathbf{W}^{(1)} + \mathbf{b}^{(1)}) \mathbf{W}^{(2)} + \mathbf{b}^{(2)})$$

message passing

- $\theta = \{\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(2)}\}$  are learnable model weights.
- Adversarial attack:** Modify node attributes  $\mathbf{X}$  and/or adjacency matrix  $\mathbf{A}$  in order to maximize classification loss
  - of an individual target node or
  - on the whole dataset/test set (global attack).

# GNN Adversarial Attacks: Challenges

1. Optimization over **discrete variables** (the graph structure). Perturbations are measured via non-convex  $L_0$  norm.
2. **Relational dependencies** between the nodes: cannot view samples in isolation.
3.  $(A', X') \approx (A, X)$ : What is a sensible measure of perturbations that do not change the semantics for (attributed) graphs?
4. **Transductive setting**: unlabeled data is **used during training**; most realistic scenario is a **poisoning attack**, where the attacker modifies the training data, which corresponds to a challenging **bilevel optimization problem**:

$$\max_{A, X} \mathcal{L}_{test}(f_{\theta^*}(A, X)) \quad s. t. \quad \theta^* = \arg \min_{\theta} \mathcal{L}_{train}(f_{\theta}(A, X))$$

# GNN Adversarial Attack: Nettack

- One of the earliest GNN attack algorithms [Zügner+ 2018].
- Targets a **single node's prediction**.

$$\mathbf{Z} = f_{\theta}(\mathbf{A}, \mathbf{X}) = \text{softmax}(\hat{\mathbf{A}} \text{ReLU}(\hat{\mathbf{A}} \mathbf{X} \mathbf{W}^{(1)}) \mathbf{W}^{(2)})$$

Linearize classifier

$$\log \mathbf{Z}' = \hat{\mathbf{A}}^2 \mathbf{X} \mathbf{W}'$$

**Structure perturbations:**  $\max_{\hat{\mathbf{A}}} \mathcal{L}'(\log \mathbf{Z}'_v)$  where  $\log \mathbf{Z}'_v = [\hat{\mathbf{A}}^2 \mathbf{c}]_v$  ← Constants

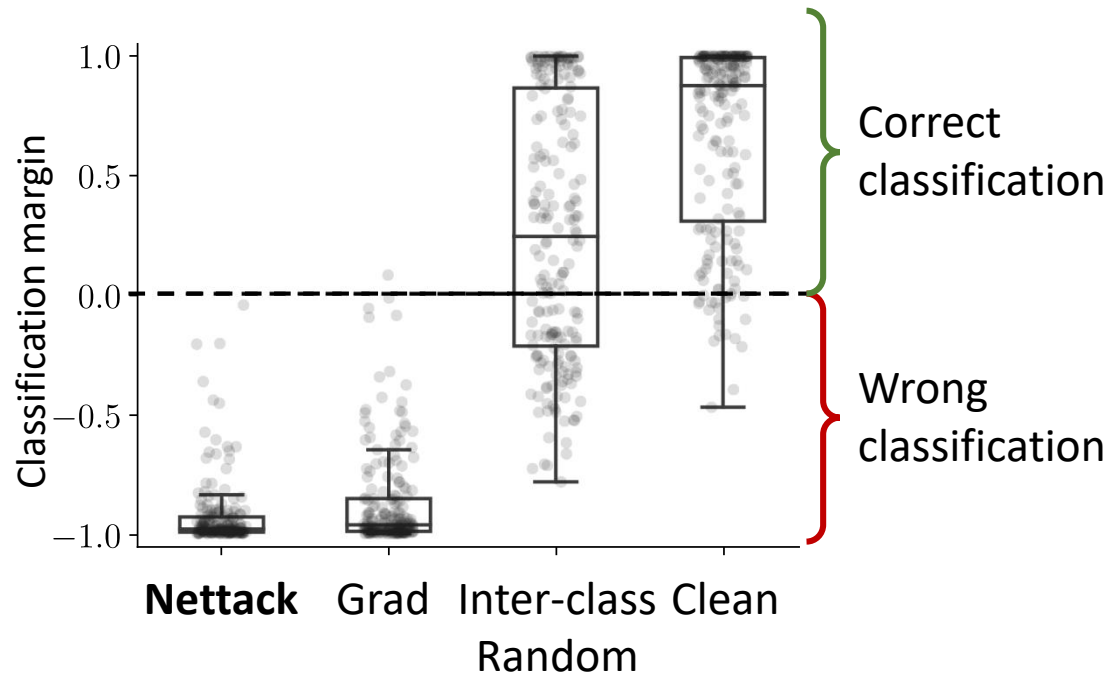
**Feature perturbations:**  $\max_{\mathbf{X}} \mathcal{L}'(\log \mathbf{Z}'_v)$  where  $\log \mathbf{Z}'_v = [\mathbf{c}_1 \mathbf{X} \mathbf{c}_2]_v$  ← Constants

→ **Greedy** pick the **optimal perturbation** at each step.

→ Uses closed-form solutions for the **optimal perturbation** at each step

# GNN Adversarial Attack: Nettack results

- Poisoning attack scenario (model is trained on perturbed data)
- Each point represents one attacked node
- Attack budget per node:  $\Delta(i) = \deg(i) + 2$



% Correct: **1.0%** 2.7% 60.8% 90.3%

# Improving Robustness

---

- GNNs are not robust under adversarial perturbations
  - specifically graph structure perturbations are harmful
- Heuristic defenses:
  - E.g. adjacency low-rank approximation via truncated Singular Value Decomposition (Entezari et al., 2020); filtering of malicious edges via attribute similarity (Wu et al., 2019)
  - However: equivalent/similar defenses for CNNs have been proven to be non-robust against worst-case perturbations
- Robust Training:
  - In form of Adversarial Training, e.g., via Projected Gradient Descent (Xu et al., 2019)
  - Or proposed together with a certification technique (upcoming topic)



## Recall: Certification (via Convex Relaxation)

**Rephrase** the original **goal**: develop an algorithm that answers the question:

“Is the GNN  $f_\theta$  around the features  $\mathbf{X}$  and adjacency matrix  $\mathbf{A}$  adversarial-free (within an  $\epsilon$ -ball(s) measured by some norm)?”

Allowed answers in the relaxed setting:

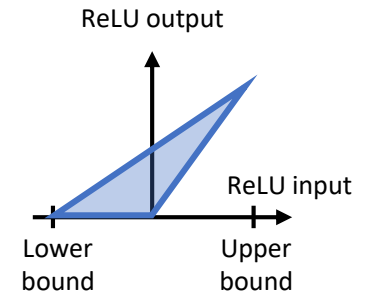
- **YES**: If for all  $\tilde{\mathbf{x}} \in \mathcal{P}_X(\mathbf{x}), \tilde{\mathbf{A}} \in \mathcal{P}_A(\mathbf{A})$ :  $\arg \max F(\tilde{\mathbf{x}}, \tilde{\mathbf{A}}) = \arg \max F(\mathbf{x}, \mathbf{A})$
- **POTENTIALLY NOT / MAYBE**: In this case we have no guarantees.
- **[NO**: If any  $\tilde{\mathbf{x}} \in \mathcal{P}_X(\mathbf{x}), \tilde{\mathbf{A}} \in \mathcal{P}_A(\mathbf{A})$ :  $\arg \max F(\tilde{\mathbf{x}}, \tilde{\mathbf{A}}) \neq \arg \max F(\mathbf{x}, \mathbf{A})$ ]



1. Graph and Attributes may change simultaneously
2. The nodes of a graph are non i.i.d.
3.  $L_0$ -ball perturbations is natural for discrete data

# Exact / Relaxed Certification

Already challenging if we are only allowed to perturb  $\mathbf{X}$



Proposed approaches so far are focusing on specific architectures and/or only attribute or structure perturbations:

- One can generalize the relaxed certification setting via linear programs to **attribute perturbations** on a **GCN** (Zügner and Günnemann, 2019).
- Certifying a **GCN** against **structure perturbations** can be formulized via a Jointly Constraint Bilinear Program (Zügner and Günnemann, 2020).
- To certify a **PPNP** model w.r.t. **structure perturbations**, we may solve a Quadratically Constrained Linear Program (Bojchevski and Günnemann, 2019).
  - under specific perturbation models (“local budget”; max  $x$  perturbations per node) one can perform certification exactly in polynomial time; for a global budget (max  $x$  perturbations overall), the problem becomes NP-hard and, thus, requires relaxation for efficiency

# Randomized Smoothing

**Recall:** Smooth classifier  $g(\mathbf{x})_c$  returns the probability that the base classifier  $f$  classifies a smoothed sample  $\tilde{\mathbf{x}} \sim \phi(\mathbf{x})$  as class  $c$

$$g(\mathbf{x})_c := \mathbb{P}(f(\phi(\mathbf{x})) = c) = \mathbb{E}_{\tilde{\mathbf{x}} \sim \phi(\mathbf{x})}(\mathbb{I}[f(\tilde{\mathbf{x}}) = c])$$

with a randomization scheme  $\phi(\mathbf{x})$ . We denote with  $c^* = \arg \max_c g(\mathbf{x})_c$  the most likely class.

**Goal:** We want to certify the smooth classifier  $g$ . That is we aim to show that for a radius  $r$  it holds:

$$\text{for all } \mathbf{x}' \text{ with } \|\mathbf{x}' - \mathbf{x}\|_0 \leq r : c^* = \arg \max_c g(\mathbf{x}')_c$$

For simplicity, we assume binary data (e.g. an unweighted graph)

→  **$L_0$ -norm radius certification.**

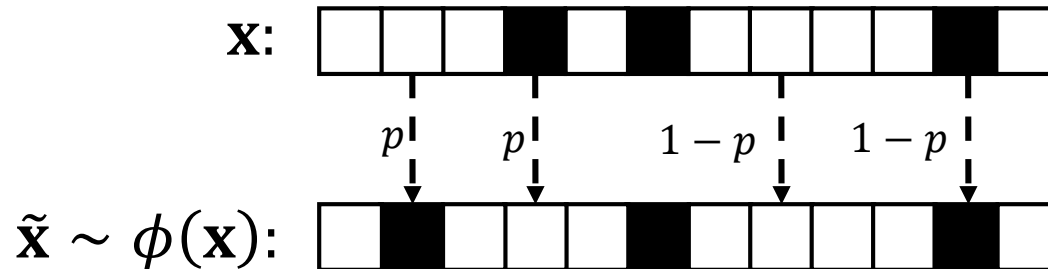
# How to Smooth Graph Data?

**Challenge:** Adding **Gaussian noise** to the **discrete graph structure** is not suitable

**Solution:** We model the  $n^2$  possible edges in the adjacency matrix as a **Bernoulli random variable**

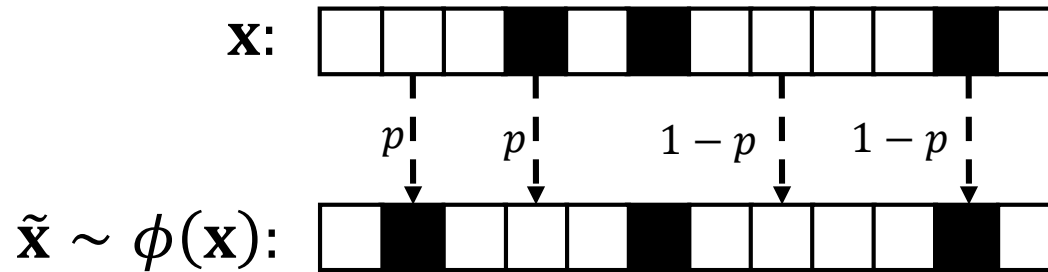
First idea: Same "flip probability" for every element

$$\tilde{\mathbf{x}} \sim \phi(\mathbf{x}) \text{ defined via } \mathbb{P}(\tilde{\mathbf{x}}_i | \mathbf{x}) = \begin{cases} p & , \tilde{\mathbf{x}}_i = 1 - \mathbf{x}_i \\ 1 - p & , \tilde{\mathbf{x}}_i = \mathbf{x}_i \end{cases}$$



# How to Smooth Graph Data?

First idea: Same "flip probability" for every element



**Problem:** Real-world graphs are typically very **sparse** ( $m \ll n^2$ ) and hence picking a meaningful  $p$  almost impossible

- Large flip probability  $p$ : most certainly we will add more random edges than original edges exist
- Small flip probability  $p$ : usually only a very few edges would be deleted

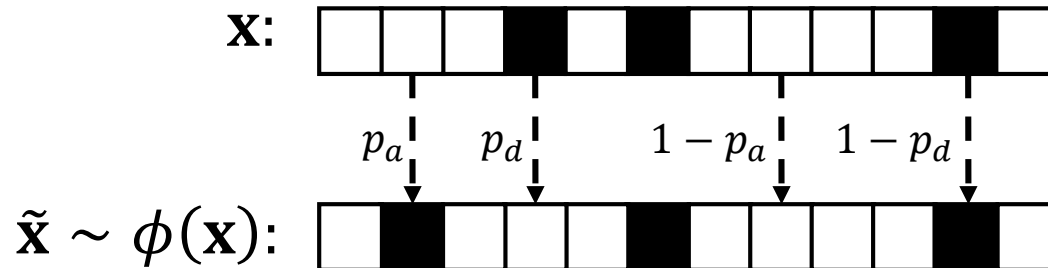
# How to Smooth Graph Data? Sparsity Matters!

**Sparsity-aware random sampling**  $\tilde{\mathbf{x}} \sim \phi(\mathbf{x})$ :

Every element of the adjacency matrix is modelled via a **Bernoulli random variable** with **data dependent probability**:

$$\mathbb{P}(\tilde{\mathbf{x}}_i | \mathbf{x}) = \begin{cases} p_d^{x_i} p_a^{1-x_i} & , \tilde{\mathbf{x}}_i = 1 - x_i \\ (1 - p_d)^{x_i} (1 - p_a)^{1-x_i} & , \tilde{\mathbf{x}}_i = x_i \end{cases}$$

That is, each of the  $n^2$  elements in the adjacency matrix is flipped with probability  $p_a$  if the value was previously 0 (no edge) or with  $p_d$  if previously an edge existed:



# Smoothed Classifier for Discrete Data

With this randomization scheme we can write:

$$\begin{aligned}
 g(\mathbf{x})_c &:= \mathbb{P}(f(\phi(\mathbf{x})) = c) = \mathbb{E}_{\tilde{\mathbf{x}} \sim \phi(\mathbf{x})}(\mathbb{I}[f(\tilde{\mathbf{x}}) = c]) \\
 &= \sum_{\tilde{\mathbf{x}}} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}) \mathbb{I}[f(\tilde{\mathbf{x}}) = c] = \sum_{\tilde{\mathbf{x}} \text{ s.t. } f(\tilde{\mathbf{x}}) = c} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}) = \sum_{\tilde{\mathbf{x}} \text{ s.t. } f(\tilde{\mathbf{x}}) = c} \prod_{i=1}^{n^2} \mathbb{P}(\tilde{\mathbf{x}}_i | \mathbf{x})
 \end{aligned}$$

We illustrate  $\mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x})$  with a hypothetical subgraph:  $\mathbf{x} =$  

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \square & \square & \blacksquare \\ \hline \end{array}\right) = (1 - p_a)^2 (1 - p_d)$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \end{array}\right) = (1 - p_a)^2 p_d$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \square & \blacksquare & \blacksquare \\ \hline \end{array}\right) = p_a (1 - p_a) (1 - p_d)$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \square & \blacksquare & \square \\ \hline \end{array}\right) = p_a (1 - p_a) p_d$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \blacksquare & \square & \blacksquare \\ \hline \end{array}\right) = p_a (1 - p_a) (1 - p_d)$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \blacksquare & \square & \square \\ \hline \end{array}\right) = p_a (1 - p_a) p_d$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \blacksquare & \blacksquare & \blacksquare \\ \hline \end{array}\right) = p_a^2 (1 - p_d)$$

$$\mathbb{P}\left(\begin{array}{|c|c|c|} \hline \blacksquare & \blacksquare & \square \\ \hline \end{array}\right) = p_a^2 p_d$$

# Worst-Case Base Classifier

Let's assume we know the value of  $g(\mathbf{x})_{c^*}$  for the original sample  $\mathbf{x}$

- Note: Of course, like in the Gaussian/continuous case, we do **not** compute this term exactly (far too expensive) but rather derive a bound based on MC samples

How far can we deviate from  $\mathbf{x}$ , e.g. obtaining  $\mathbf{x}'$ , and still **guarantee** that we do not change the prediction, i.e. still have  $\arg \max_c g(\mathbf{x})_c = c^* = \arg \max_c g(\mathbf{x}')_c$ ?

→ Similarly to the Gaussian/continuous randomized smoothing, to answer this question, we can inspect the **worst-case base classifier**.

Since the worst-case base classifier has a simple form (e.g. linear in the Gaussian case), once we know it, it is “rather simple” to obtain the certification radius



# The Space of Base Classifiers

- Recall: we only assumed knowledge about the value of  $g(\mathbf{x})_{c^*}$ 
  - We do **not** know the output of the actual base classifier  $f$  at every possible input
- Various base classifiers  $h$  fulfill the property  $\sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}})=c^*} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}) = g(\mathbf{x})_{c^*}$ 
  - Let  $\mathcal{H}$  denote the set of all these base classifiers; clearly  $f \in \mathcal{H}$

$h_1 \in \mathcal{H}$		$h_2 \in \mathcal{H}$	
$\mathbb{P}(\begin{array}{ c c c } \hline \square & \square & \blacksquare \\ \hline \end{array})$	$= (1 - p_a)^2(1 - p_d) =$	$\mathbb{P}(\begin{array}{ c c c } \hline \square & \square & \blacksquare \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \square & \square & \square \\ \hline \end{array})$	$= (1 - p_a)^2 p_d =$	$\mathbb{P}(\begin{array}{ c c c } \hline \square & \square & \square \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \square \\ \hline \end{array})$	$= p_a(1 - p_a)(1 - p_d) =$	$\mathbb{P}(\begin{array}{ c c c } \hline \square & \blacksquare & \square \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \square & \blacksquare & \blacksquare \\ \hline \end{array})$	$= p_a(1 - p_a)p_d =$	$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \square \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \square & \blacksquare \\ \hline \end{array})$	$= p_a(1 - p_a)(1 - p_d) =$	$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \square & \blacksquare \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \square & \square \\ \hline \end{array})$	$= p_a(1 - p_a)p_d =$	$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \square & \square \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \square \\ \hline \end{array})$	$= p_a^2(1 - p_d) =$	$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \square \\ \hline \end{array})$	
$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \blacksquare \\ \hline \end{array})$	$= p_a^2 p_d =$	$\mathbb{P}(\begin{array}{ c c c } \hline \blacksquare & \blacksquare & \blacksquare \\ \hline \end{array})$	

# How to Find the Worst-Case Base Classifier?

For any **chosen location**  $\mathbf{x}'$ , we can express the **worst-case base classifier** as a minimization problem:

$$\min_{h \in \mathcal{H}} \sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}}) = c^*} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}') \quad \left( \leq \sum_{\tilde{\mathbf{x}} \text{ s.t. } f(\tilde{\mathbf{x}}) = c} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}') = g(\mathbf{x}')_{c^*} \right),$$

## Intuition:

We search for a base classifier  $h$  such that the resulting smooth classifier

- maintains the probability mass for the clean sample  $\mathbf{x}$ ,  
i.e.  $\mathbb{P}(h(\phi(\mathbf{x})) = c^*) = g(\mathbf{x})_{c^*}$  //  $h \in \mathcal{H}$
- and simultaneously minimizes the probability mass at the  
perturbed sample  $\mathbf{x}'$ , i.e.  $\mathbb{P}(h(\phi(\mathbf{x}')) = c^*) = \sum_{\tilde{\mathbf{x}} \text{ s.t. } h(\tilde{\mathbf{x}}) = c^*} \mathbb{P}(\tilde{\mathbf{x}} | \mathbf{x}')$

# Solution for the Worst-Case Base Classifier

The previous minimization problem can be formulated as a linear program!

Denote with  $\tilde{\mathbf{x}}^{(i)}$  the enumeration of all possible  $\tilde{\mathbf{x}}$ .

$$\begin{aligned} & \min_{\mathbf{h}} \sum_i \mathbf{h}_i \mathbb{P}(\tilde{\mathbf{x}}^{(i)} | \mathbf{x}') \\ & \text{subject to } \sum_i \mathbf{h}_i \mathbb{P}(\tilde{\mathbf{x}}^{(i)} | \mathbf{x}) = g(\mathbf{x})_{c^*} \text{ and } 0 \leq \mathbf{h}_i \leq 1 \end{aligned}$$

The vector  $\mathbf{h}$  represents the worst-case base classifier:  $\mathbf{h}_i$  indicates whether  $h(\tilde{\mathbf{x}}^{(i)})$  outputs  $c^*$  ( $\mathbf{h}_i = 1$ ) or some other class

- Technically it is a soft classifier (like logistic regression) since  $0 \leq \mathbf{h}_i \leq 1$

# Solution for the Worst-Case Base Classifier

Interesting fact: This is a very special LP, which can efficiently and exactly be solved with a greedy approach

- Initialize all  $\mathbf{h}_i$  with zero
  - Compute likelihood ratios  $\eta_i = \frac{\mathbb{P}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x})}{\mathbb{P}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x}' )}$  and sort them
    - i.e. get indices  $j_1, j_2, j_3, \dots$  such that  $\eta_{j_1} \geq \eta_{j_2} \geq \eta_{j_3} \geq \dots$
  - For  $k = 1, \dots$  set  $\mathbf{h}_{j_k} = 1$  while budget  $\sum_i \mathbf{h}_i \mathbb{P}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x}) = g(\mathbf{x})_{c^*}$  is not used up
    - i.e. process the sorted indices from left to right and assign a 1 (again: at the “switch point” we might have  $0 < \mathbf{h}_j < 1$  to consume the budget fully).
- Result: We do not even have to solve an optimization problem! We just sort based on the **likelihood ratio** and assign class  $c^*$  to the “left part”
- Note the similarity to the linear classifier in the Gaussian/continuous case
- The worst-case base classifier has a very simple form

## Some Details We Skip

- Knowing the worst-case base classifier, enables us to find the certification radius  $r$  (technically we even have two radii: addition/deletion)
  - Core insight: The general form of the worst-case classifier is always the same, independent of which  $\mathbf{x}'$  we consider; similar to the Gaussian/continuous case
- In the linear program the dimensionality of  $\mathbf{h}$  would be enormous (all possible graphs). We can use the fact that only the likelihood ratio  $\eta_i = \frac{\mathbb{P}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x})}{\mathbb{P}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x}')}$  matters for the solution.
  - Intuition: Group together all graphs that have the same value for  $\eta_i$  into one large region → dimensionality of  $\mathbf{h}$  equals to the number of regions
  - Indeed we have only a small number of regions: linear in the radius/dimensionality of the input; very fast certification possible
- For further details we refer to (Bojchevski et al., 2020).

**Most importantly, this randomized smoothing technique works for all models with binary input data: GNNs, CNNs, SVMs, Decision Trees, ...**

# Questions

---

1. Is a projected-gradient-descent (PGD) attack on a GNN via the graph structure a good idea? Why or why not?
2. Suppose you want to determine the worst-case structure perturbation  $\Delta$ , which is limited to (i) insert or (ii) remove at most  $k$  edges. How many possible perturbations are there (in big-O notation w.r.t. the number of nodes  $N$  and number of edges  $E$ )?
3. Given a graph with 2810 nodes and 7336 edges. What value of  $p_a$  do we need to choose if in expectation we want to sample 7336 further edges?

# Summary

---

- GNNs are **not robust** to adversarial attacks.
- GNN robustness/certification is a highly **active research area**.
  - To date there exists **no defense** against **structure attacks** that consistently improves results; standard methods such as **adversarial training** do not seem to work well.
- **Robustness certification** of GNNs is challenging but possible
  - specialized approaches enable to exploit structure of the GNN models
- **Randomized smoothing** can be adapted to **discrete input** data via **Bernoulli random variables**
  - We draw Monte Carlo samples for  $g(\mathbf{x})$  and obtain the certified radii analytically.
  - Most importantly, **randomized smoothing with the proposed noise model works for all models with binary input data**: GNNs, CNNs, SVMs, Decision Trees, ...

## References: Attacks

---

- Bojchevski, Aleksandar, and Stephan Günnemann. "Adversarial attacks on node embeddings via graph poisoning." ICML 2019.
- Zügner, Daniel, Amir Akbarnejad, and Stephan Günnemann. "Adversarial attacks on neural networks for graph data." *KDD* 2018.
- Zügner, Daniel and Stephan Günnemann. "Adversarial attacks on Graph Neural Networks via Meta Learning." *ICLR 2019*.



## References: Improving Robustness

---

- Negin Entezari, Saba A. Al-Sayouri, Amirali Darvishzadeh, and Evangelos E. Papalexakis. All you need is Low (rank): Defending against adversarial attacks on graphs. WSDM 2020 - Proceedings of the 13th International Conference on Web Search and Data Mining, pages 169–177, 2020.
- Huijun Wu, Chen Wang, Yuriy Tyshetskiy, Andrew Docherty, Kai Lu, and Liming Zhu. Adversarial examples for graph data: Deep insights into attack and defense. IJCAI International Joint Conference on Artificial Intelligence, 2019.
- Kaidi Xu, Hongge Chen, Sijia Liu, Pin Yu Chen, Tsui Wei Weng, Mingyi Hong, and Xue Lin. Topology attack and defense for graph neural networks: An optimization perspective. IJCAI International Joint Conference on Artificial Intelligence, 2019.

## References: Certification

---

- Daniel Zügner and Stephan Günnemann. Certifiable robustness and robust training for graph convolutional networks. Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pages 246–256, 2019.
- Daniel Zügner and Stephan Günnemann. Certifiable Robustness of Graph Convolutional Networks under Structure Perturbations. Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2020.
- Aleksandar Bojchevski and Stephan Günnemann. Certifiable Robustness to Graph Perturbations. (NeurIPS), 2019
- Aleksandar Bojchevski, Johannes Klicpera and Stephan Günnemann. Efficient Robustness Certificates for Discrete Data: Sparsity-Aware Randomized Smoothing for Graphs, Images and More, ICML 2020